

AWS Cloud Technology and Services

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS



Alex Kuntz

Head of Cloud Curriculum, DataCamp

Chapter 1: Compute Services



AWS EC2



AWS Lambda

Chapter 2: Databases, networking, and storage



AWS DynamoDB



AWS S3

Chapter 3: AI, machine learning, and more

Artificial Intelligence services



Amazon Translate



Amazon Polly



Amazon Lex



Amazon Comprehend



Amazon Forecast



Amazon Rekognition



Amazon CodeGuru

Prerequisites

- No technical experience needed
- A basic understanding of AWS could help

INTERACTIVE COURSE

Introduction to AWS

[Continue](#)[Bookmark](#)

• Beginner

🕒 2 hours

▶ 12 videos

⌂ 39 exercises

👤 10,584 participants

2650 XP

Course format

aws

Services

Search

[Option+S]

N. Virginia

datacamp-learner-user @ 3397-1279-7442

Console Home

Reset to default layout

Add widgets

Recently visited

S3

AWS Billing Conductor

Billing and Cost Management

IAM

Lambda

EC2

Certificate Manager

IAM Identity Center

Secrets Manager

Trusted Advisor

Security Hub

CloudWatch

View all services

Applications (0)

Create application

Region: US East (N. Virginia)

us-east-1 (Current Region)

Find applications

< 1 >

Name	Description	Region	Originating account
No applications Get started by creating an application.			
Create application			

Go to myApplications

Welcome to AWS

AWS Health

Cost and usage

Let's practice!

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS

Amazon Elastic Compute Cloud

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS



Alex Kuntz


Head of Cloud Curriculum, DataCamp

What is Amazon Elastic Compute Cloud (EC2)?

- Provides resizable compute capacity in the cloud
- Each individual EC2 machine is referred to as an instance

Key characteristics:

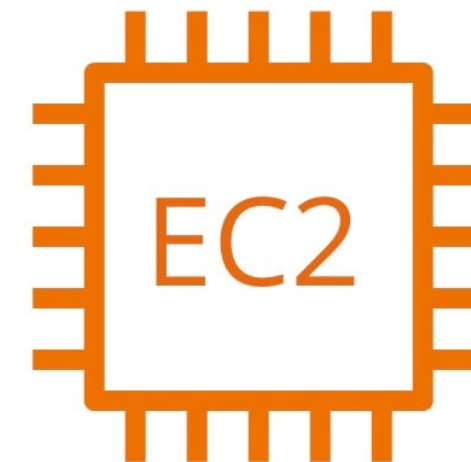
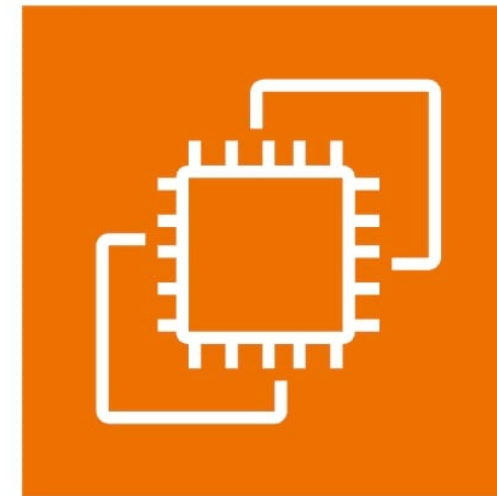
- Ability to scale up or down based on demand
- Availability of a varied range of instance types for specialized use cases



EC2 instance types

AWS offers six categories of EC2 instances for specialized workloads


- General purpose
- Compute optimized
- Memory optimized
- Storage optimized
- Accelerated computing
- High Performance Computing (HPC) optimized



General purpose and storage optimized instances

General purpose instances

- Balance of compute, memory, and networking resources
- Use cases:
 - Hosting dynamic websites
 - Maintaining code repositories



Storage optimized instances


- High, sequential read and write access to large datasets
- Use cases:
 - Data warehousing
 - Refactoring large relational databases



Compute and memory optimized instances


Compute optimized instances

- Compute-intensive and high-performance workloads
- Use cases:
 - Scientific simulations
 - Financial modeling



Memory optimized instances

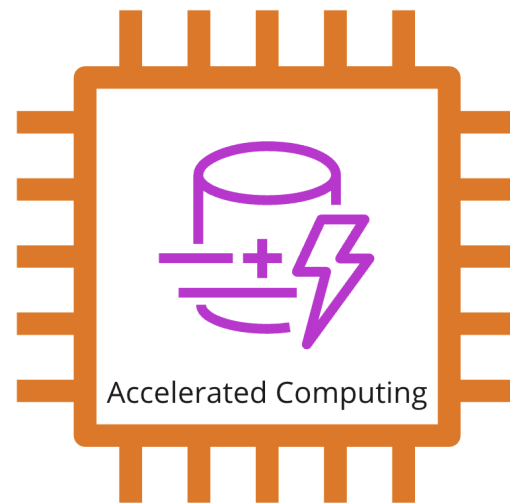
- Memory-intensive workloads not requiring high storage
- Use cases:
 - Real-time stream data analytics
 - Generating close captions



Specialized compute instances


Accelerated computing instances

- Contain specialized hardware accelerators, like GPUs or FPGAs
- Use cases:
 - Deep learning
 - Rendering gaming graphics

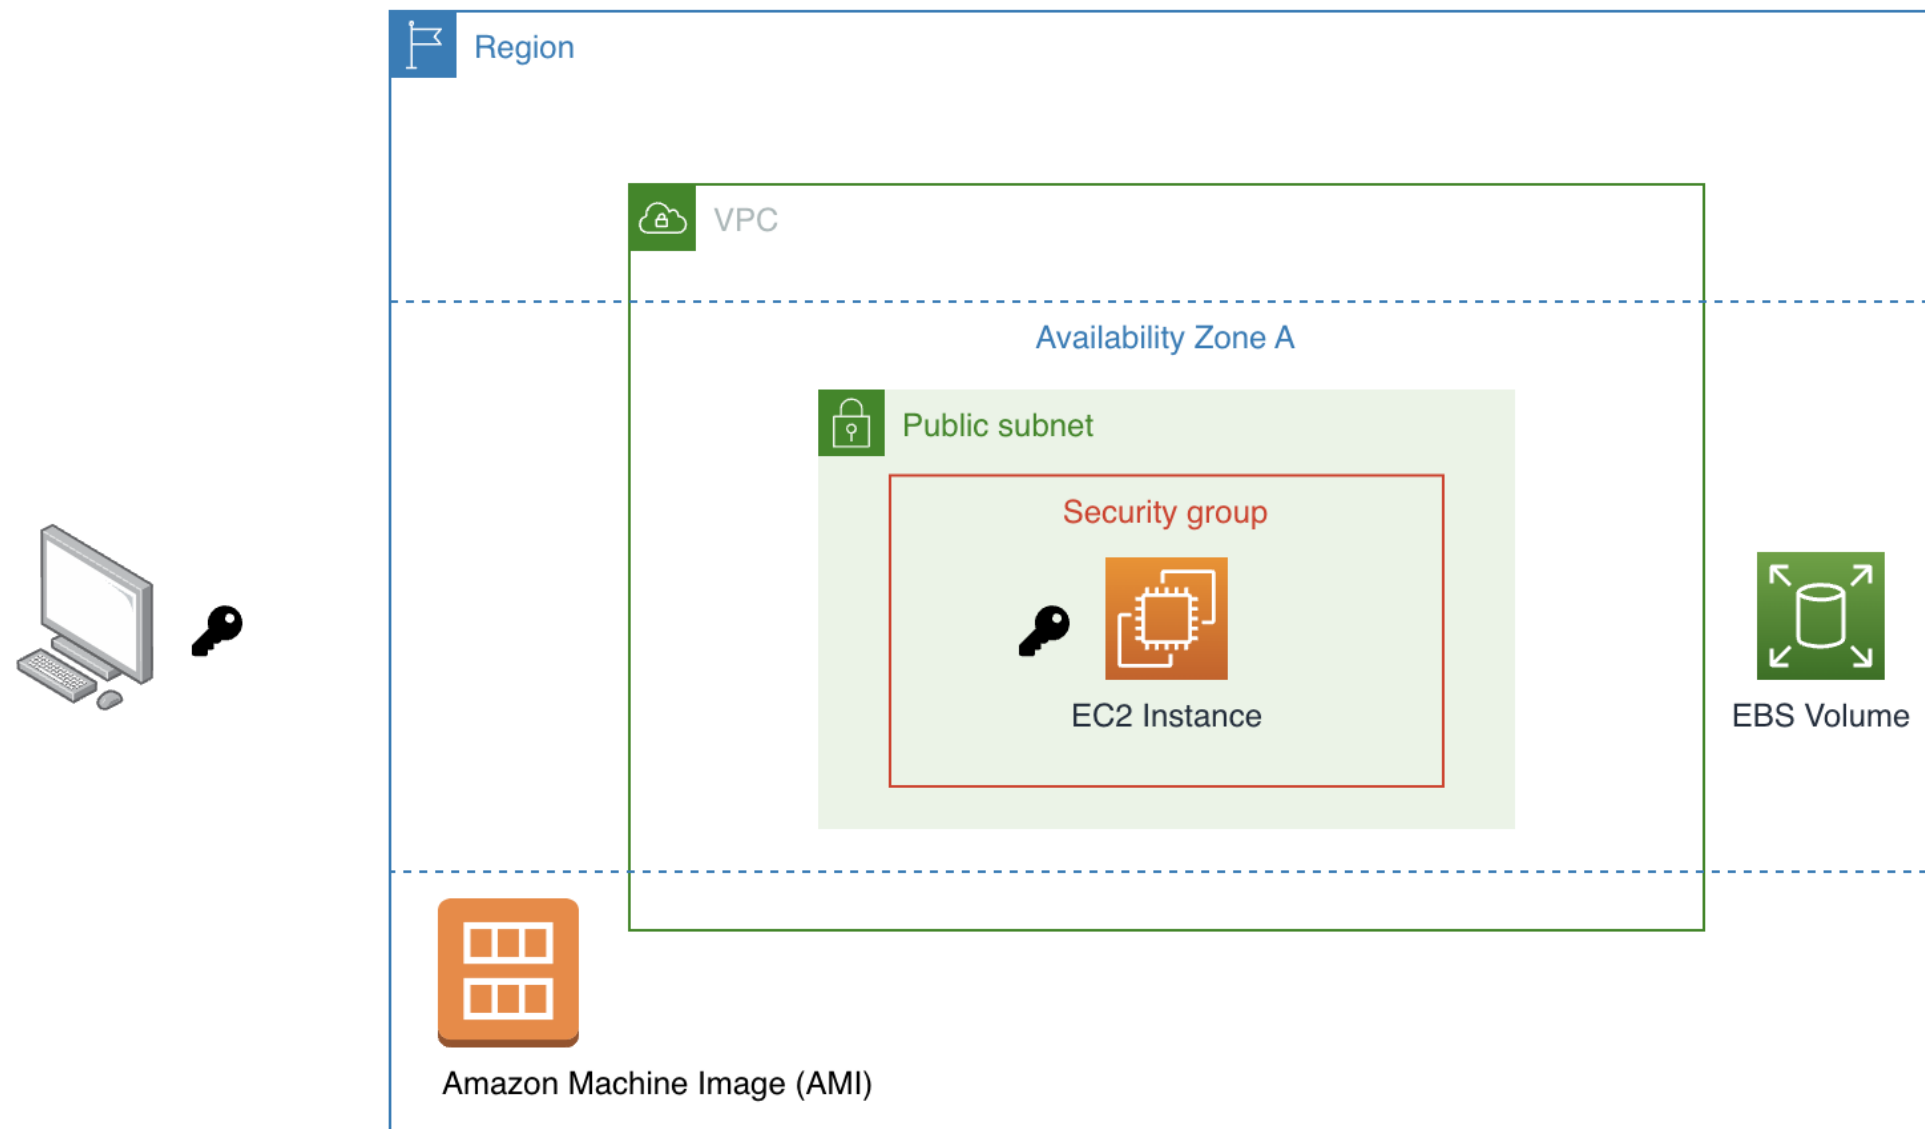


HPC optimized instances

- Best price performance for running high performance workloads at scale
- Use cases:
 - Weather forecasting
 - Crash simulations



Creating your EC2 instance



Connecting to your EC2 instance: SSH Client

- SSH connects with a private key
- Those keys must be managed

Connect to instance [Info](#)

Connect to your Instance `i-XXXXXXXXXXXX` (MyWebServer) using any of these options

[EC2 Instance Connect](#) | [Session Manager](#) | [SSH client](#) | [EC2 serial console](#)

Instance ID

`i-XXXXXXXXXXXX` (MyWebServer)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is `Test.pem`
3. Run this command, if necessary, to ensure your key is not publicly viewable.
`chmod 400 ".pem"`
4. Connect to your Instance using its Public DNS:
`ec2-XXXXXXXXXXXX.compute-1.amazonaws.com`

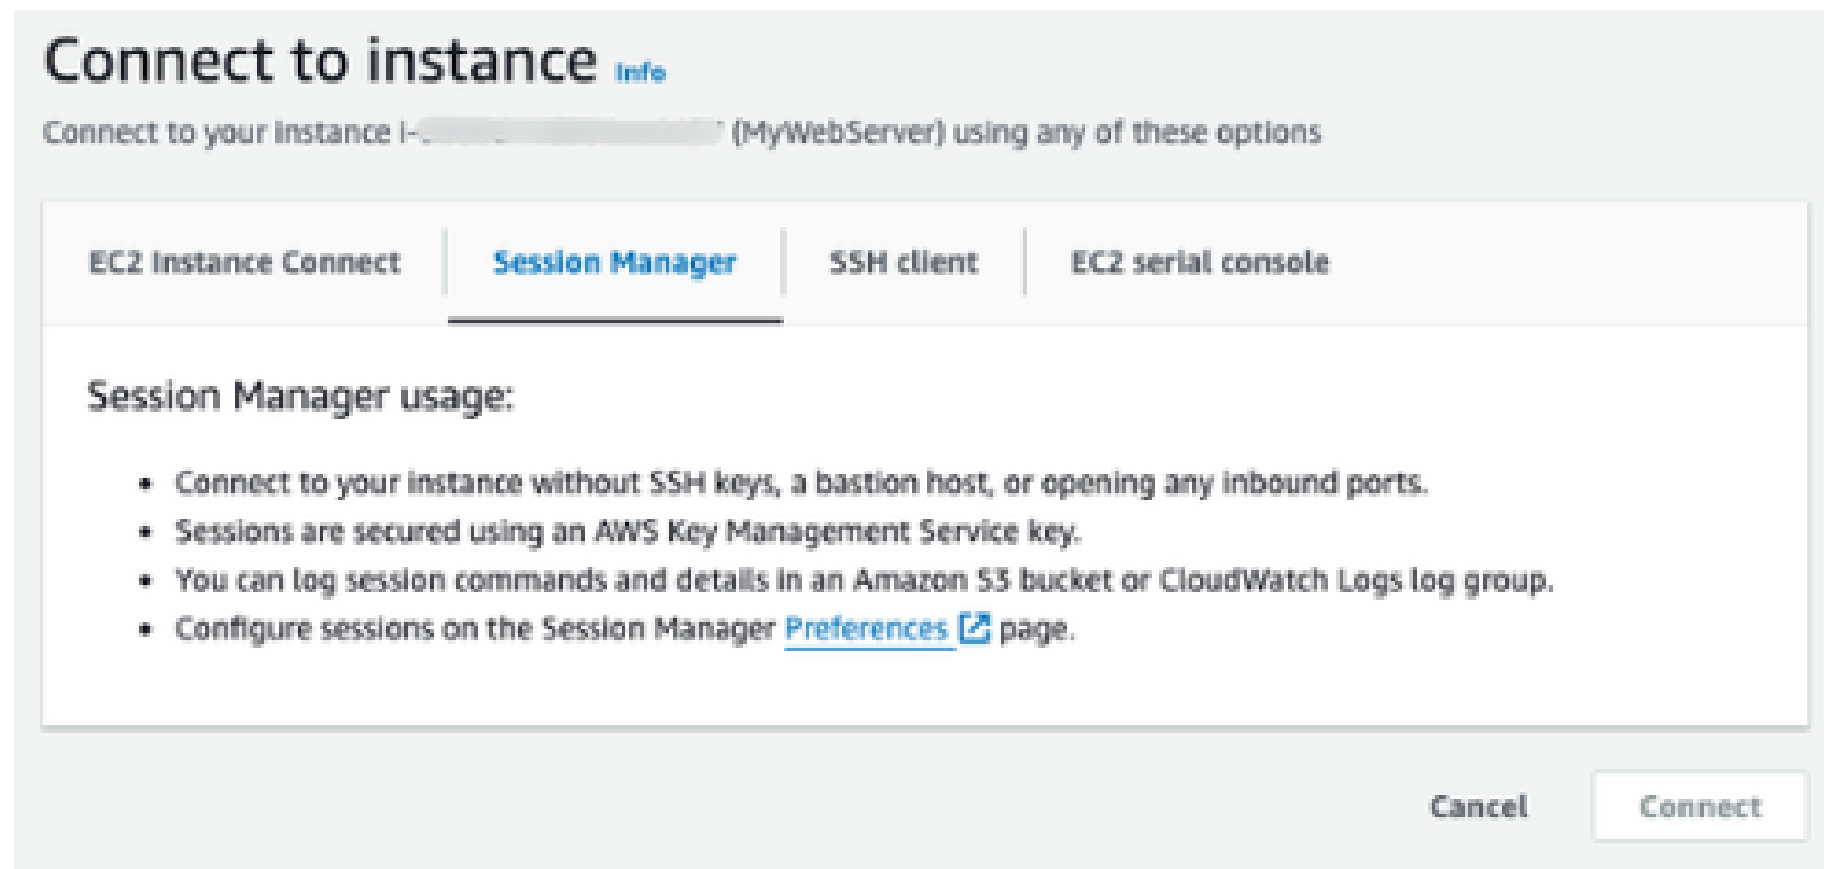
Example:

```
ssh -i ".pem" ec2-user@ec2-XXXXXXXXXXXX.compute-1.amazonaws.com
```

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

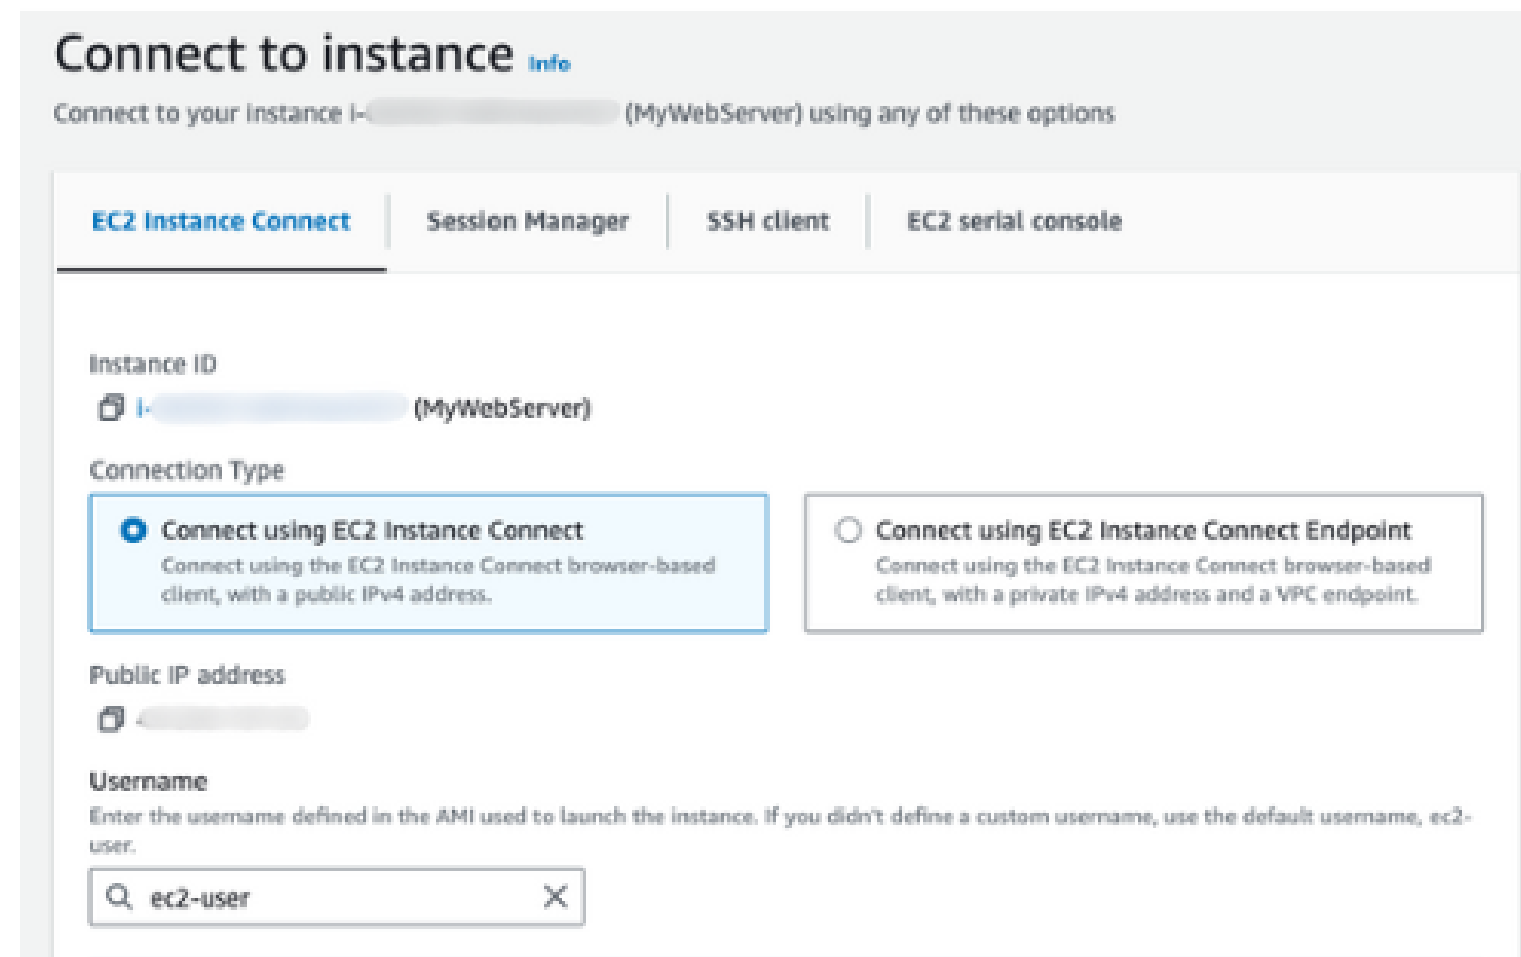
Connecting to your EC2 instance: AWS Session Manager

- Keyless access via the Management Console
- Integrates with Identity and Access Management (IAM)



Connecting to your EC2 instance: EC2 Instance Connect

- Browser-based connection
- Quick and temporary access



Connect to instance [Info](#)

Connect to your instance **i-12345678** (MyWebServer) using any of these options

EC2 Instance Connect | Session Manager | SSH client | EC2 serial console

Instance ID
i-12345678 (MyWebServer)

Connection Type

☒ **Connect using EC2 Instance Connect**
Connect using the EC2 Instance Connect browser-based client, with a public IPv4 address.

☐ **Connect using EC2 Instance Connect Endpoint**
Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint.

Public IP address
44.192.164.123

Username
Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

Let's practice!

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS

Load Balancing and Auto-scaling

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS



Alex Kuntz

Head of Cloud Curriculum, DataCamp

Load balancing in AWS

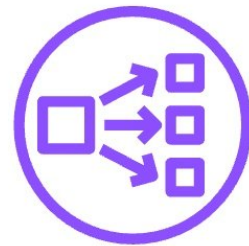
Load balancing ensures even distribution of incoming traffic among multiple EC2 instances, preventing overload on a single server

- Ensures high availability
- Provides horizontal scaling

Types of load balancers in AWS



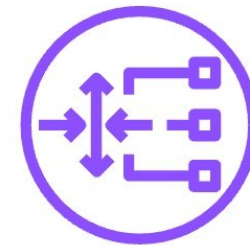
Classic Load Balancer



Network Load Balancer



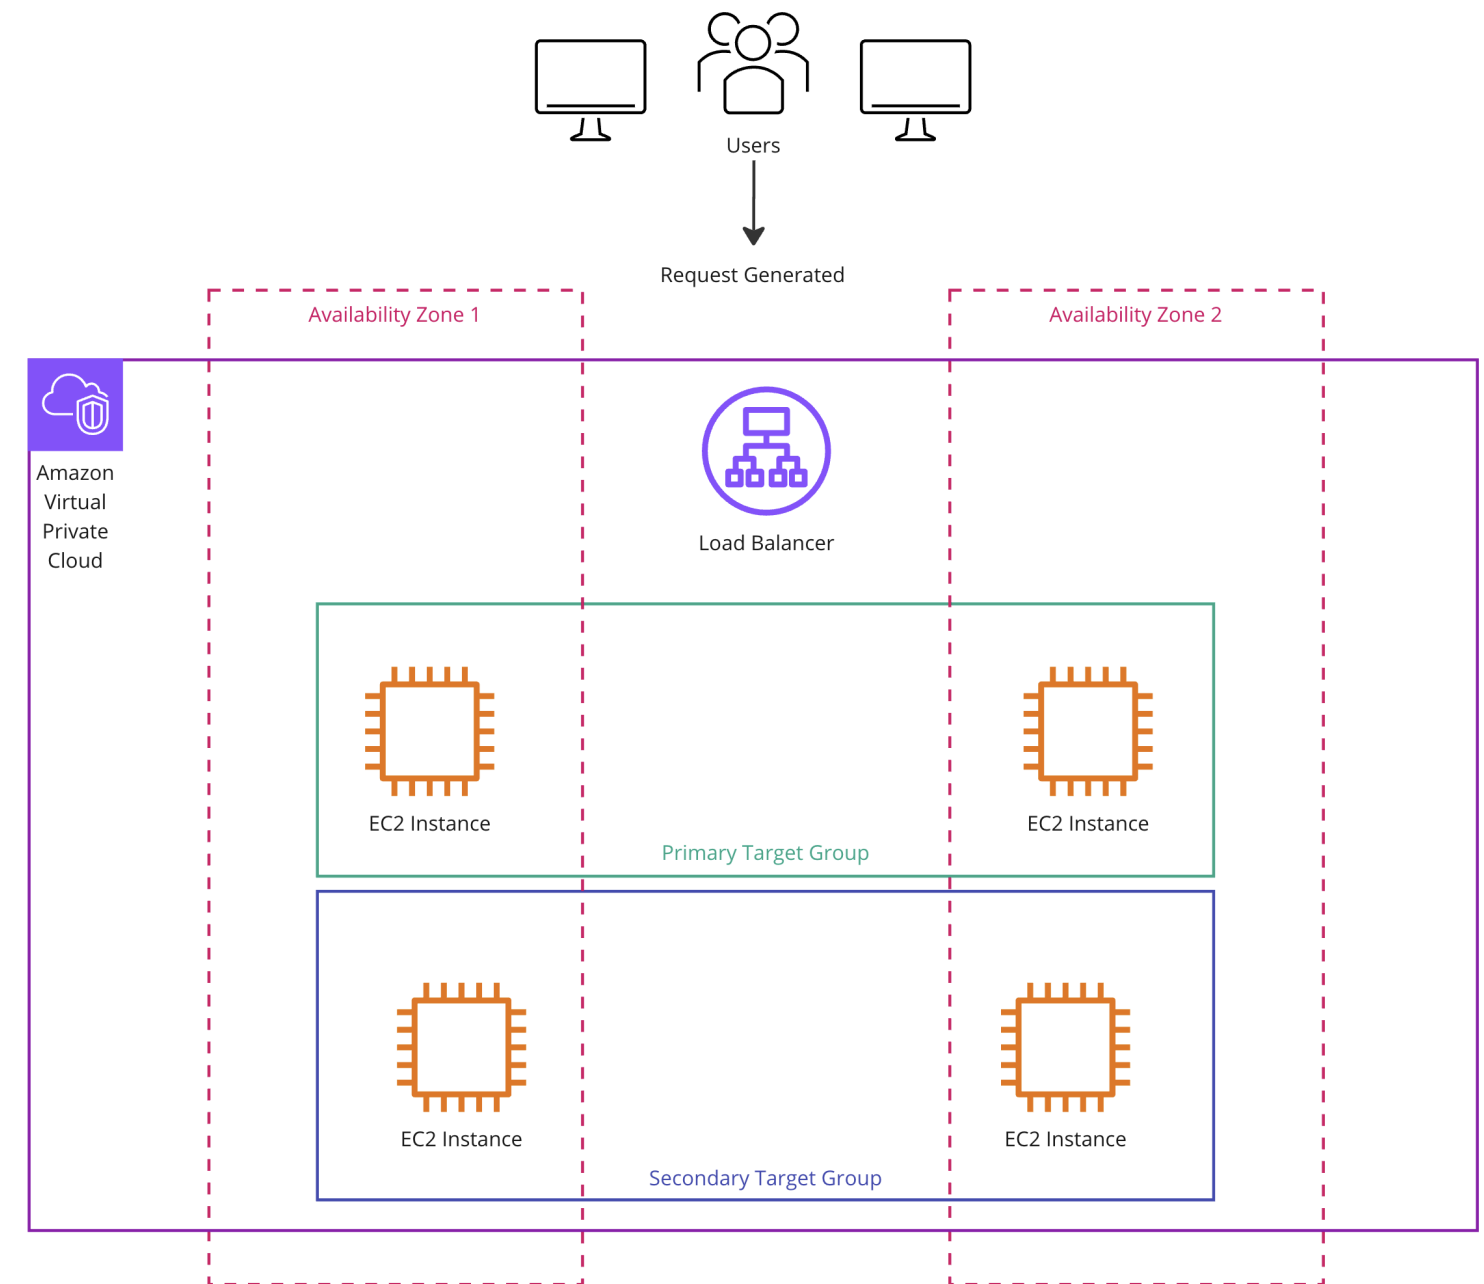
Application Load Balancer



Gateway Load Balancer

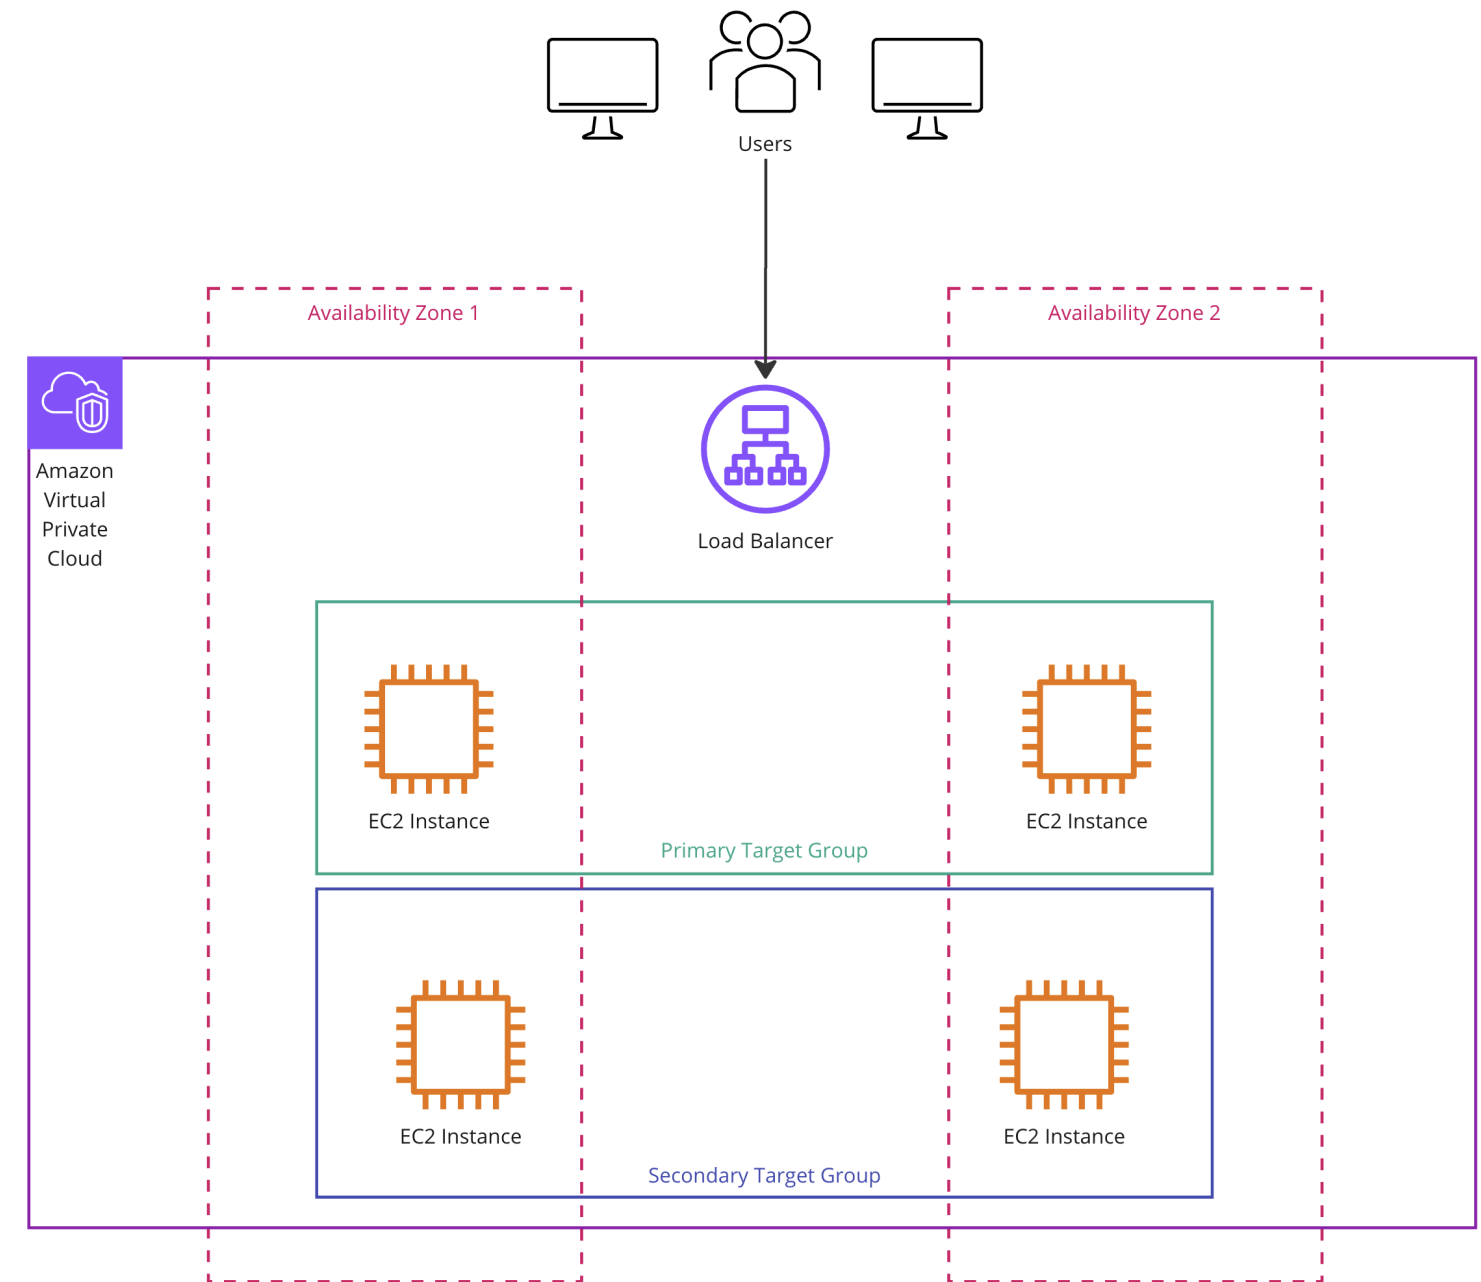
How does load balancing work?

1. Users send requests



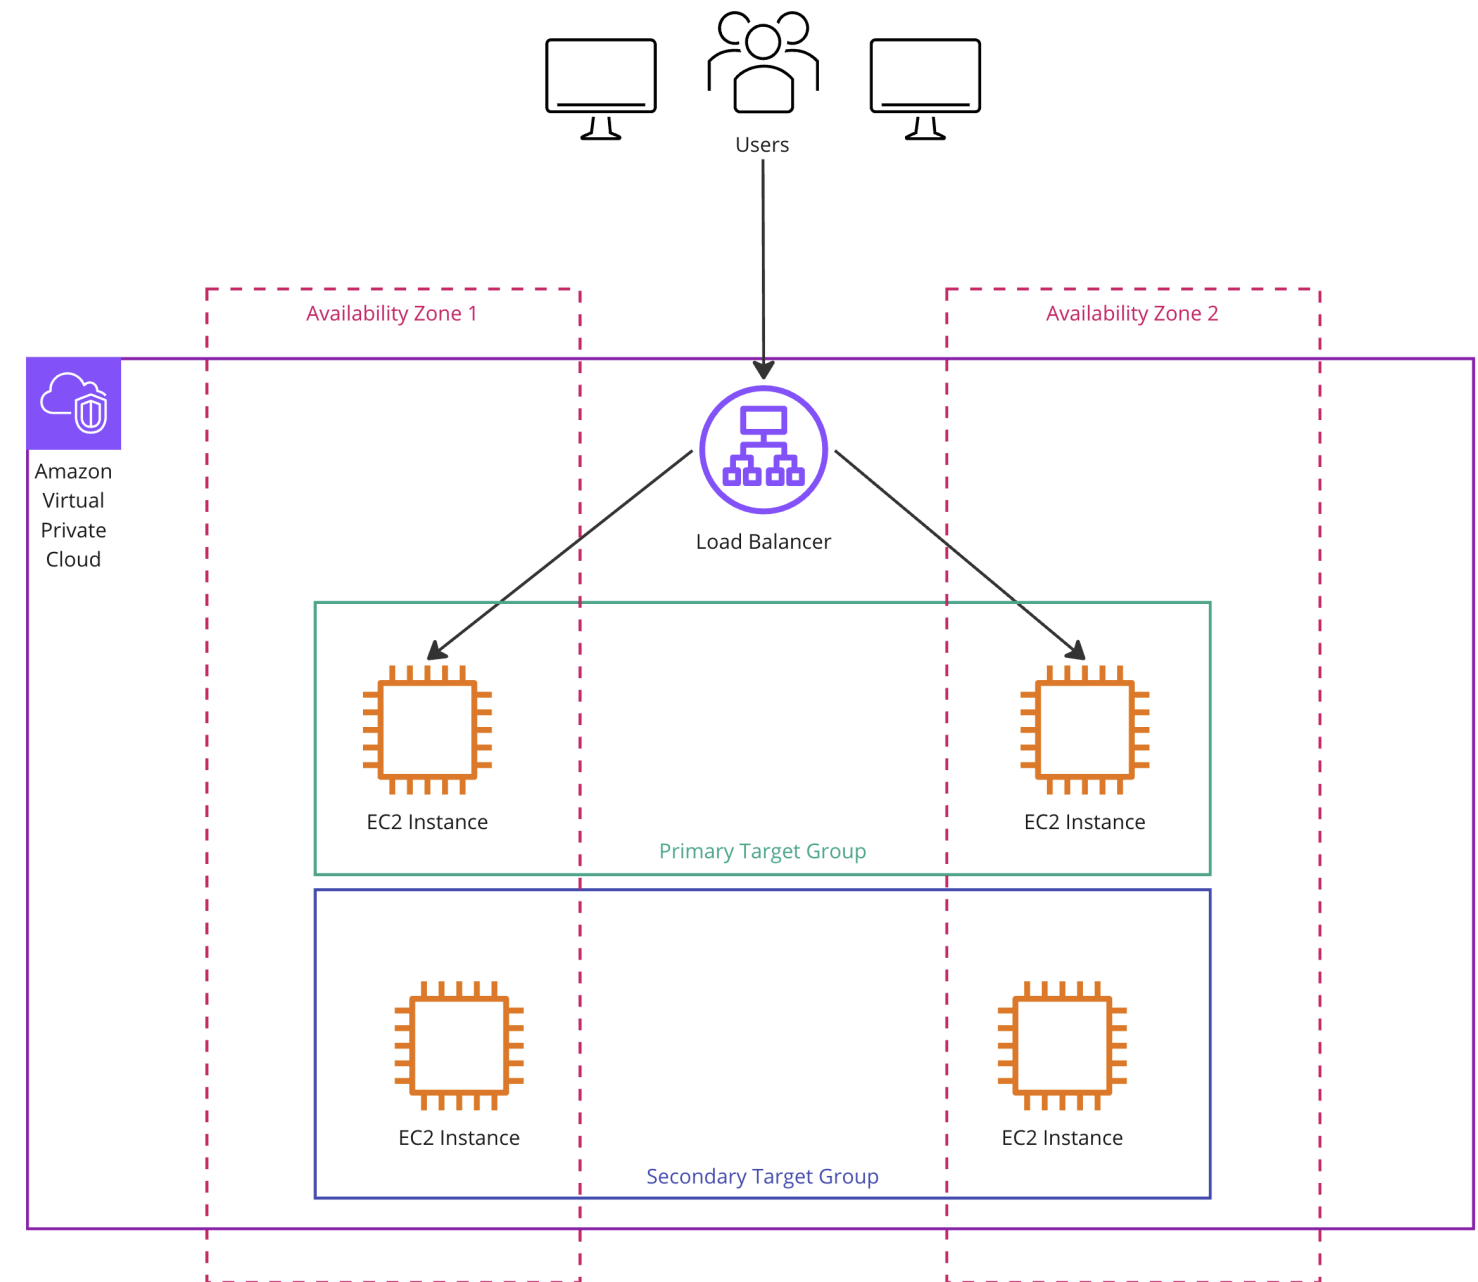
How does load balancing work?

1. Users send requests
2. Requests hit the load balancer



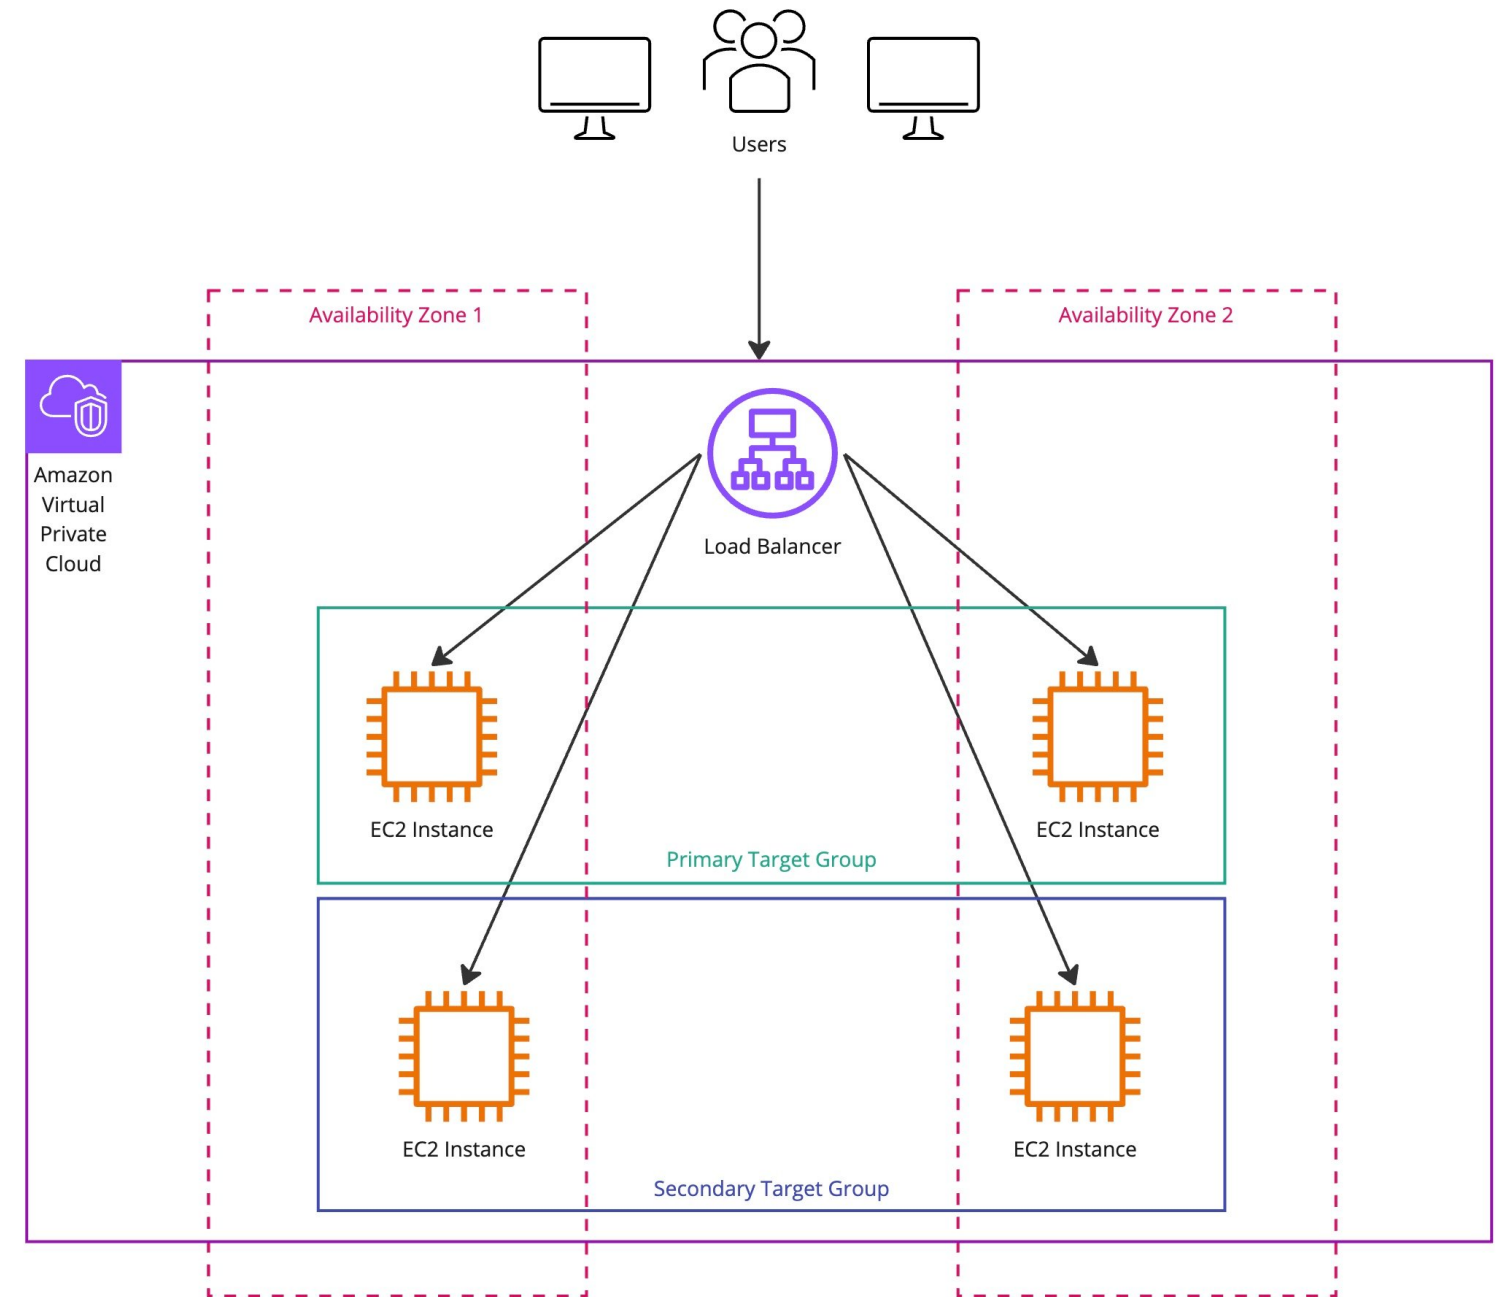
How does load balancing work?

1. Users send requests
2. Requests hit the load balancer
3. Primary target group is instantiated by the application load balancer first



How does load balancing work?

1. Users send requests
2. Requests hit the load balancer
3. Primary target group is instantiated by the application load balancer first
4. If demand increases, the load balancer activates the secondary target group and distributes the load across all instances



What is compute elasticity?

Elasticity ensures your system can scale up or down based on demand, providing flexibility in resource allocation

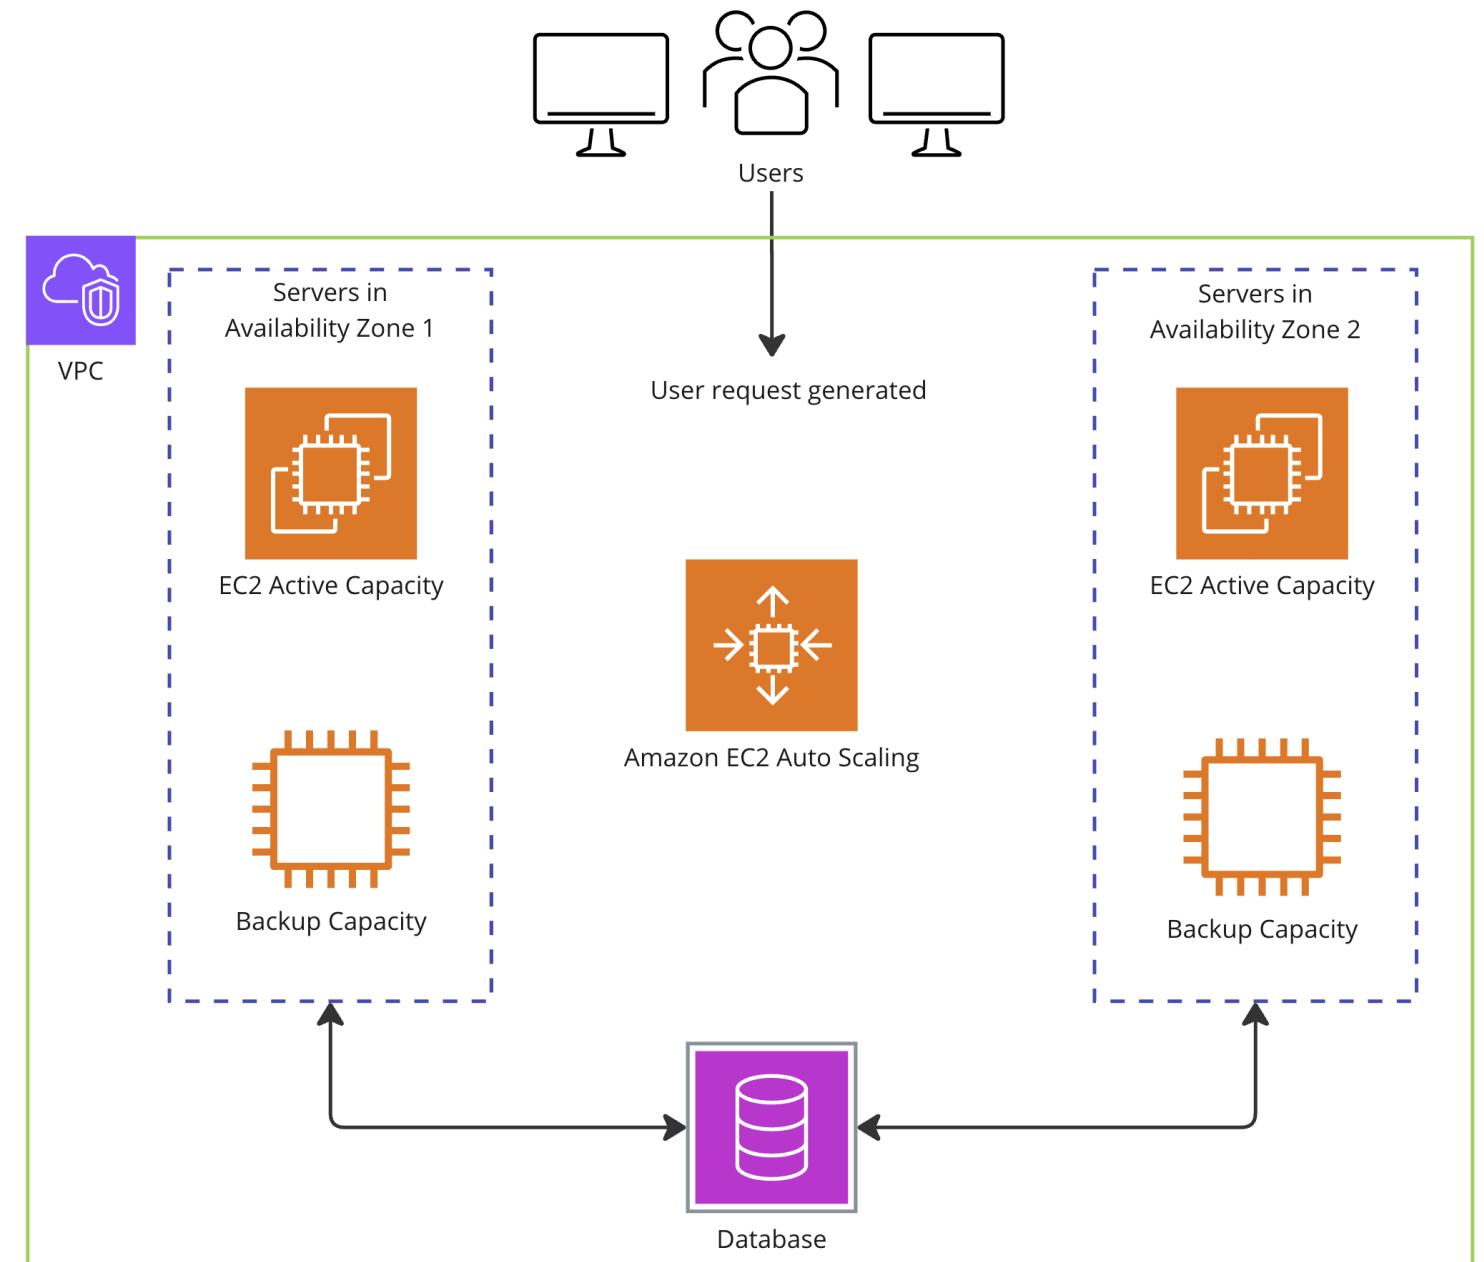
- EC2 instances achieve elasticity through EC2 Auto Scaling

What is EC2 Auto Scaling?

- Automatically adjust the number of active instances based on usage and requirement
- Optimize costs
- Prevent over-provisioning

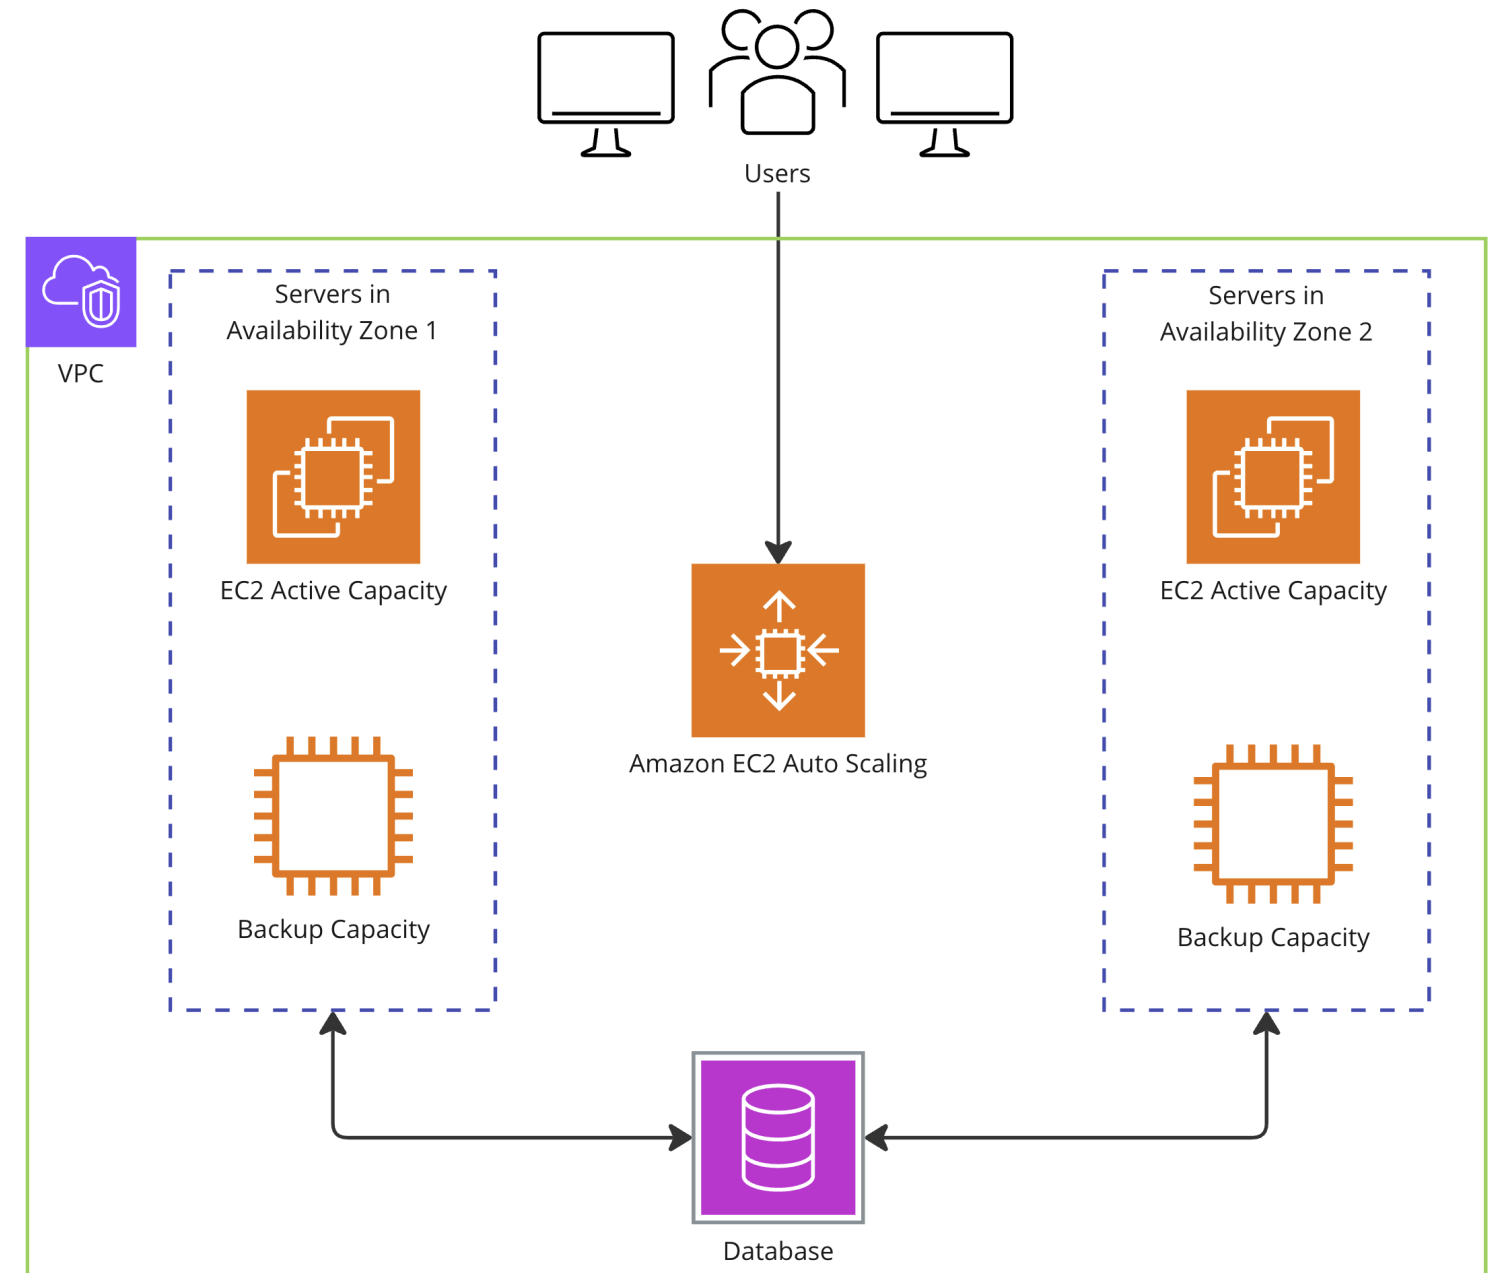
How does auto-scaling work?

1. Users send requests



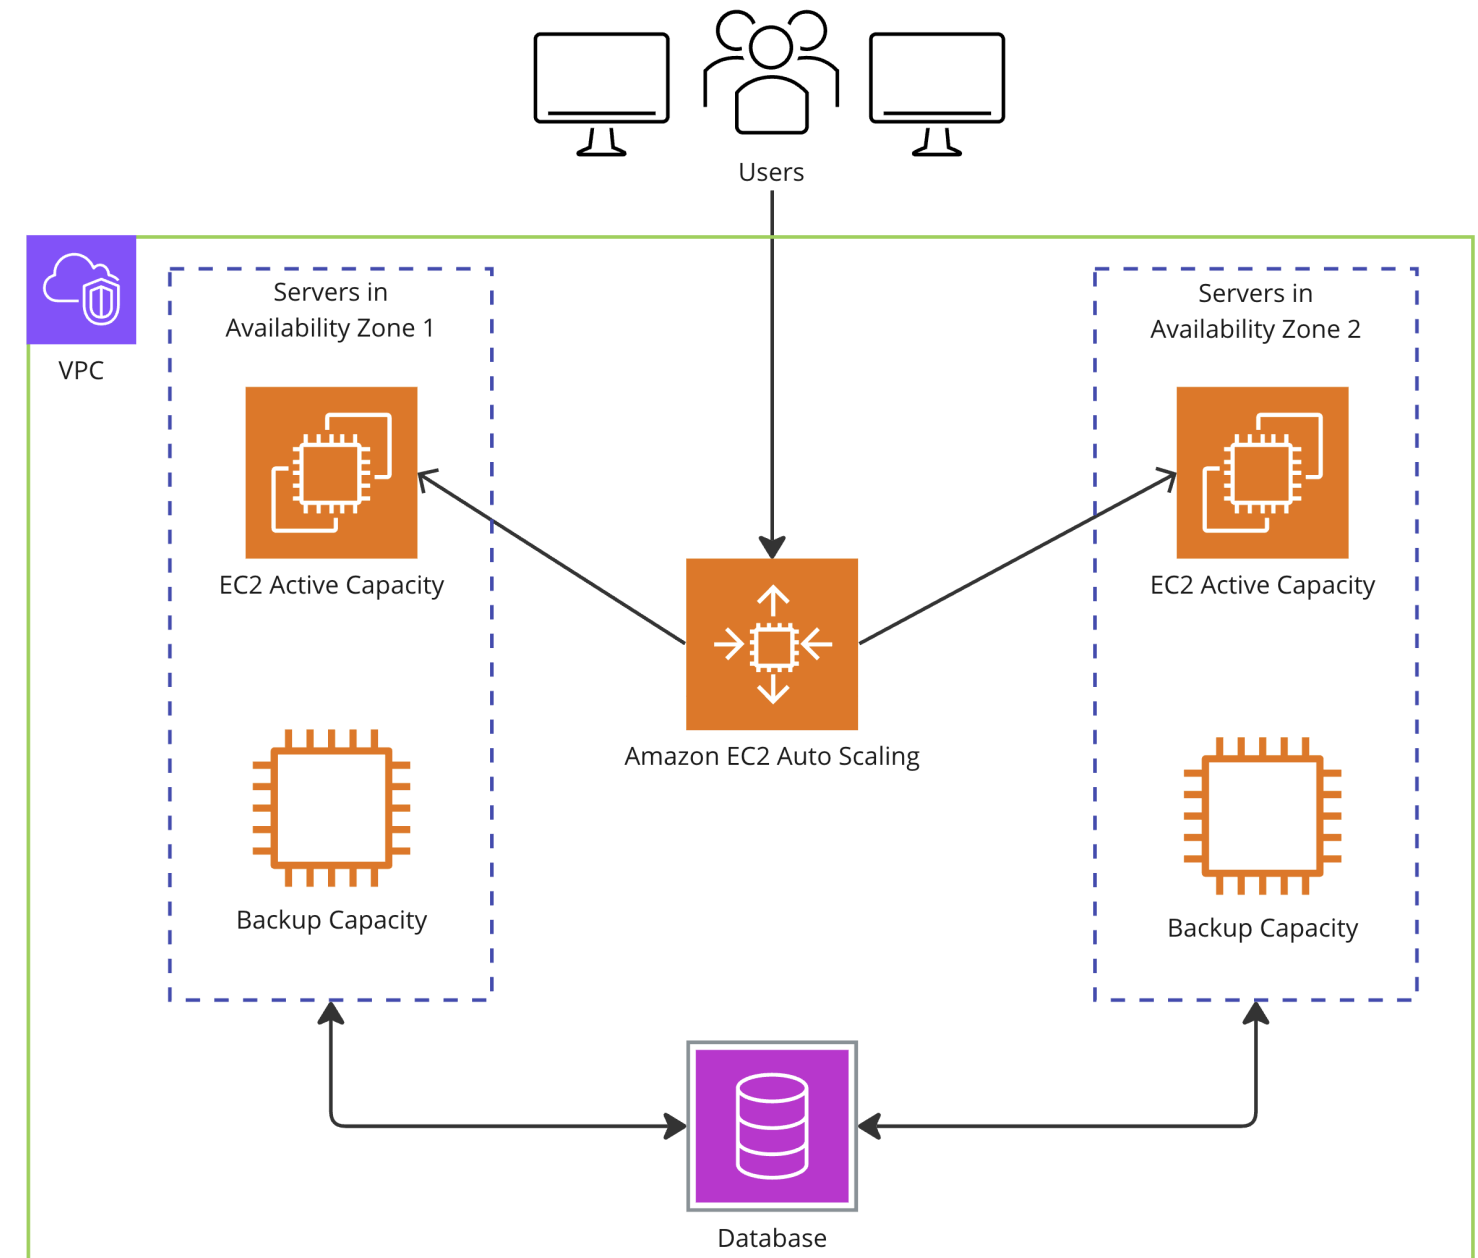
How does auto-scaling work?

1. Users send requests
2. The requests are routed to EC2 Auto Scaling service



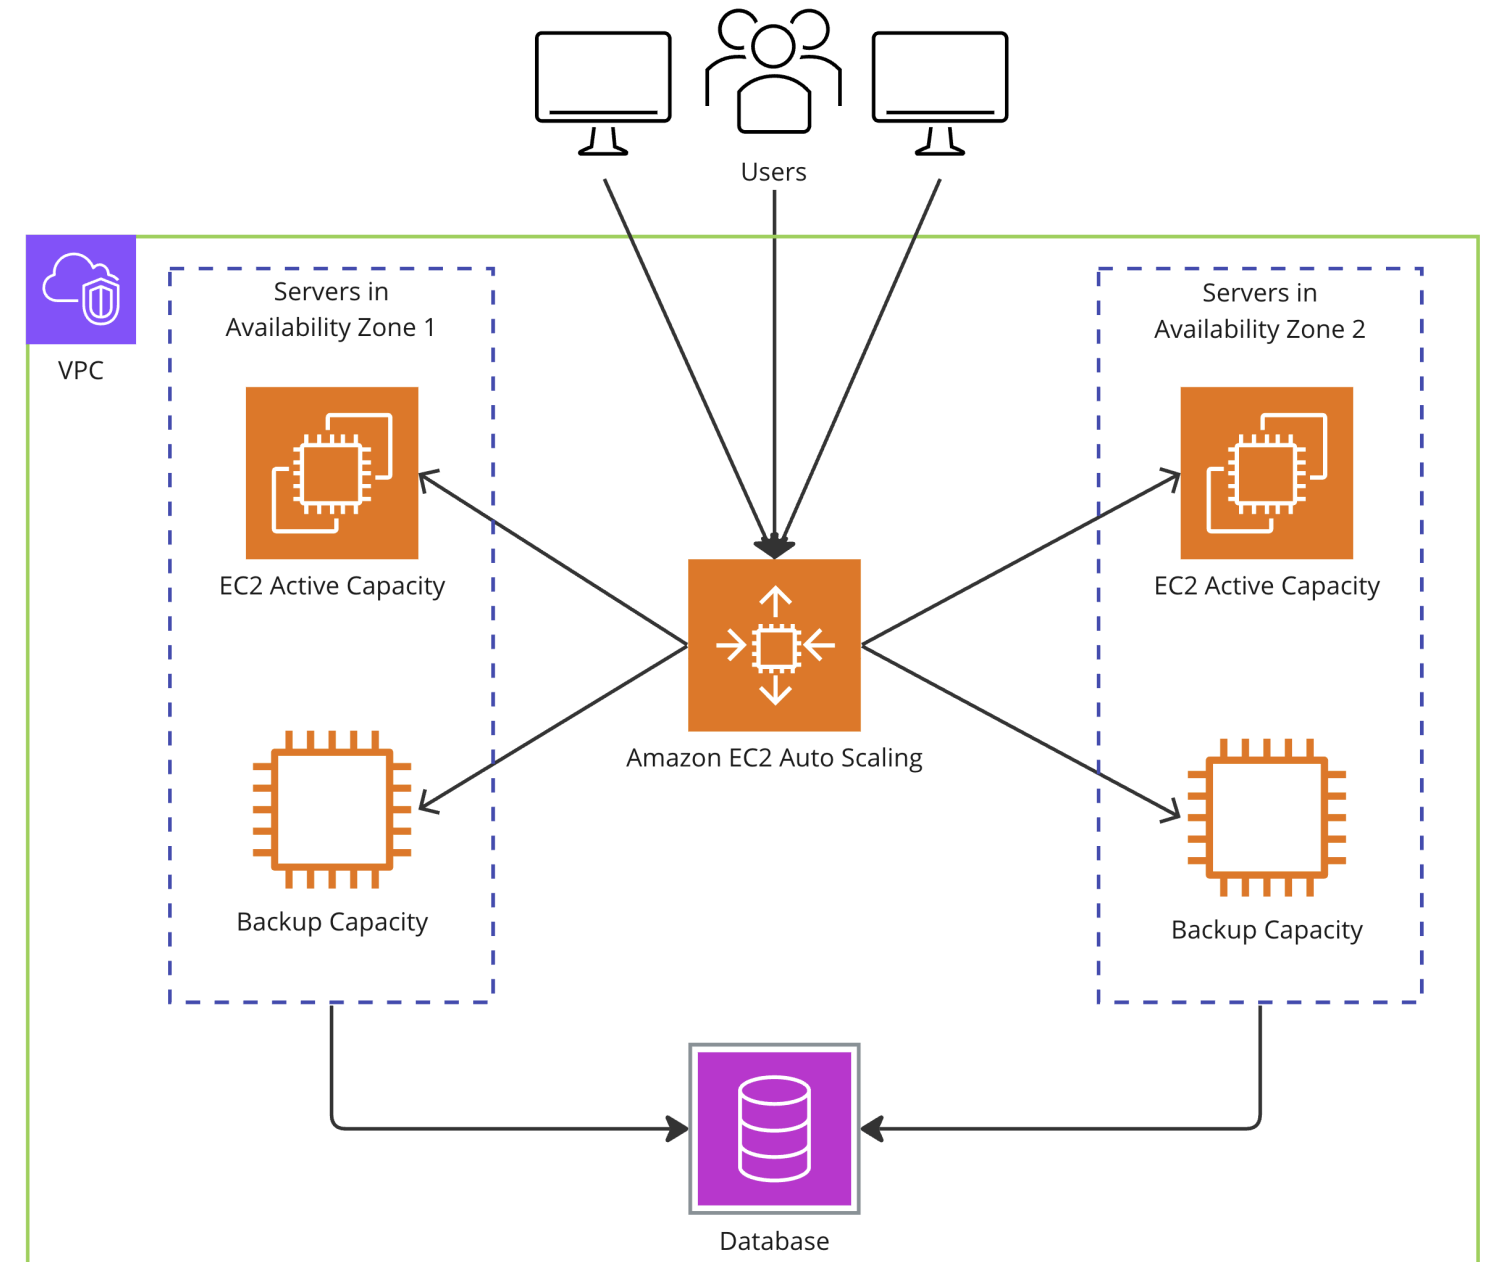
How does auto-scaling work?

1. Users send requests
2. The requests are routed to EC2 Auto Scaling service
3. The service then routes requests to the active EC2 instances



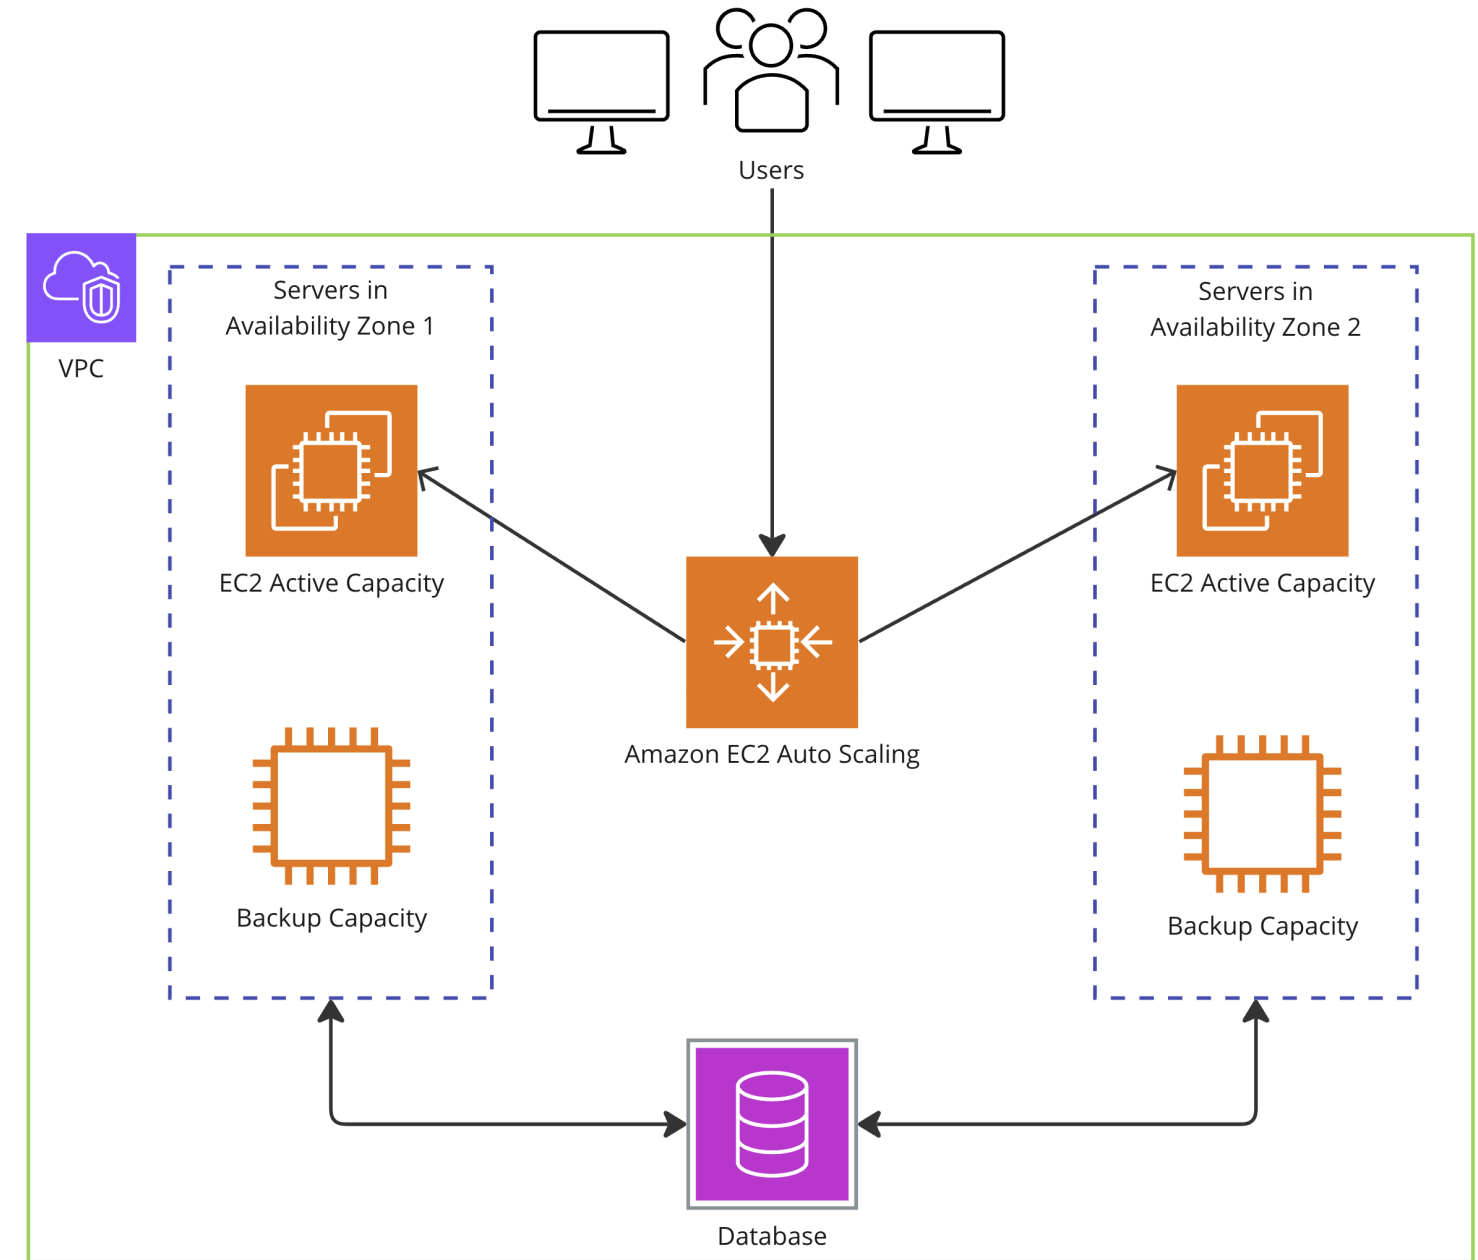
How does auto-scaling work?

1. Users send requests
2. The requests are routed to EC2 Auto Scaling service
3. The service then routes requests to the active EC2 instances
4. If demand increases, it starts adding new EC2 instances to manage the additional load



How does auto-scaling work?

1. Users send requests
2. The requests are routed to EC2 Auto Scaling service
3. The service then routes requests to the active EC2 instances
4. If demand increases, it starts adding new EC2 instances to manage the additional load
5. As demand goes down, the newly added EC2 instances are shut down



Load balancing vs. auto-scaling

Load balancing

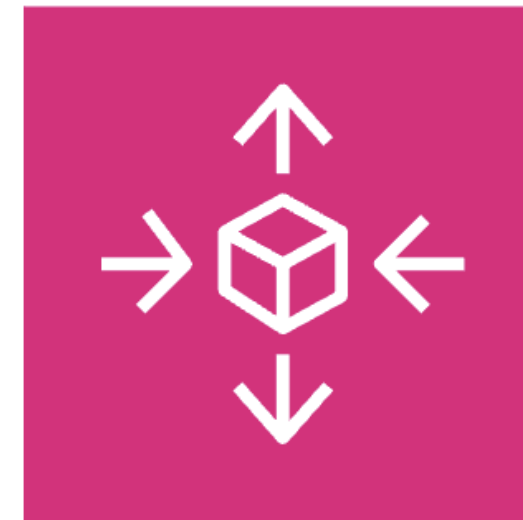
- Route traffic evenly
- Utilize existing EC2 instances



AWS Load Balancer

Auto-scaling

- Ensure demand is always met
- Ability to add/remove EC2 instances



EC2 Auto Scaling

Let's practice!

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS

Serverless Compute

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS



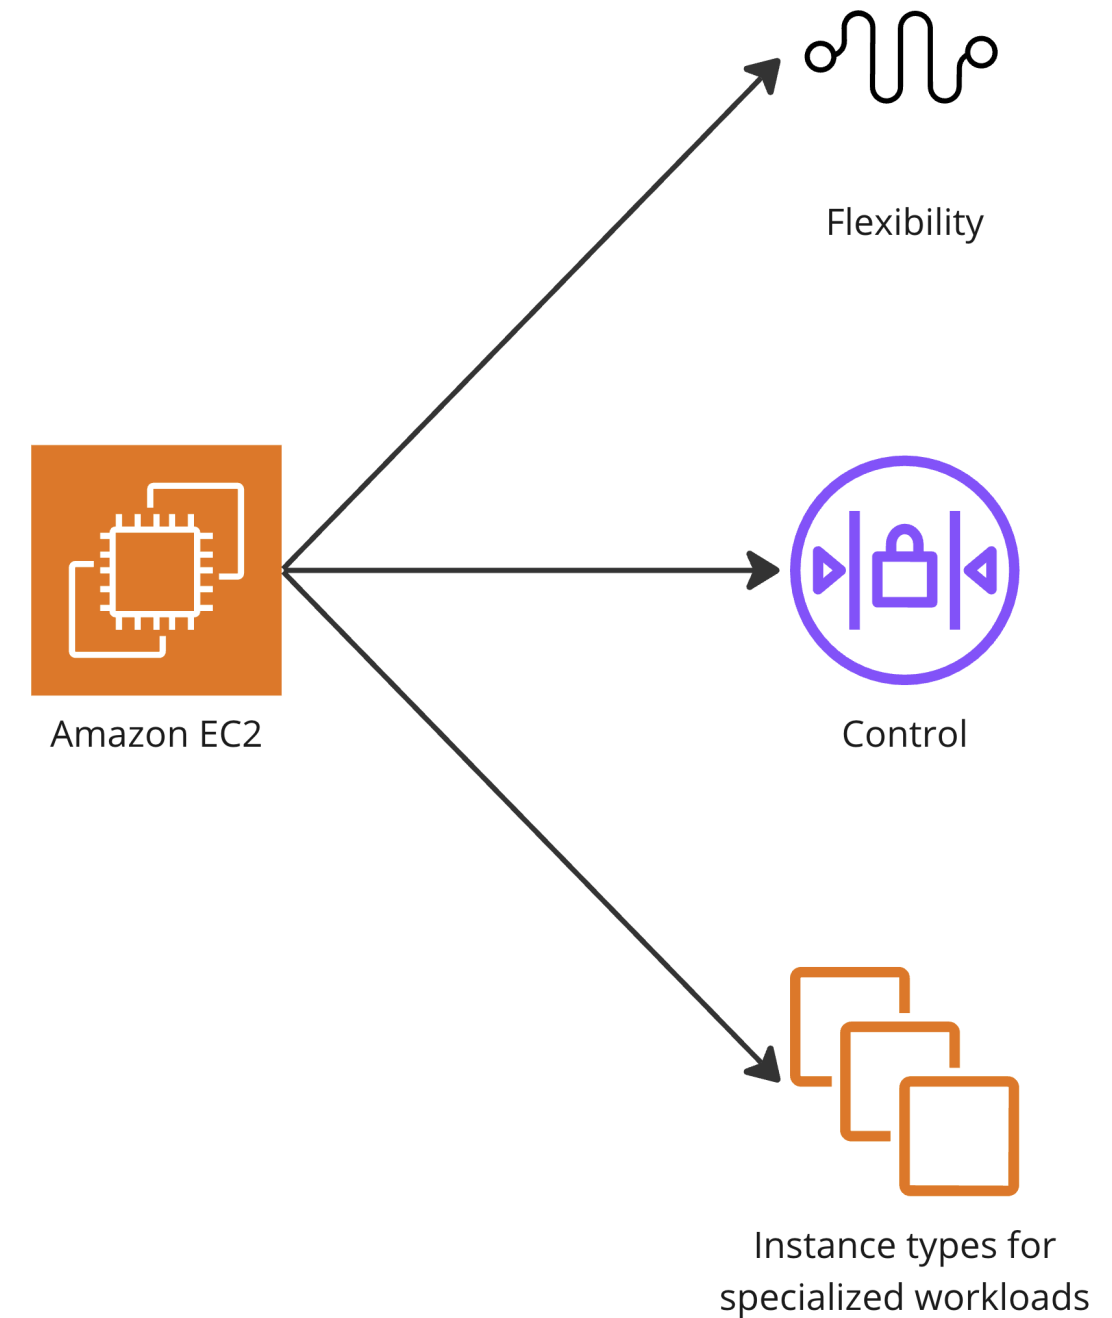
Alex Kuntz

Head of Cloud Curriculum, DataCamp

Recap traditional compute

EC2 Recap:

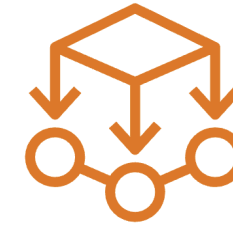
- Amazon EC2 is a service that provides compute capacity in the AWS cloud
- Using EC2 gives higher flexibility and control
- Variety of EC2 instance types optimized for different workloads



Evolving needs: beyond traditional compute

Today's Demands:

- Need for modular, microservices architectures




Microservice architecture


Evolving needs: beyond traditional compute

Today's Demands:

- Need for modular, microservices architecture
- Rapid scaling capabilities to meet fluctuating demands



Microservice architecture




Rapid scaling


Evolving needs: beyond traditional compute

Today's Demands:

- Need for modular, microservices architecture
- Rapid scaling capabilities to meet fluctuating demands
- Automated infrastructure management setup without manual interventions



Microservice architecture



Rapid scaling

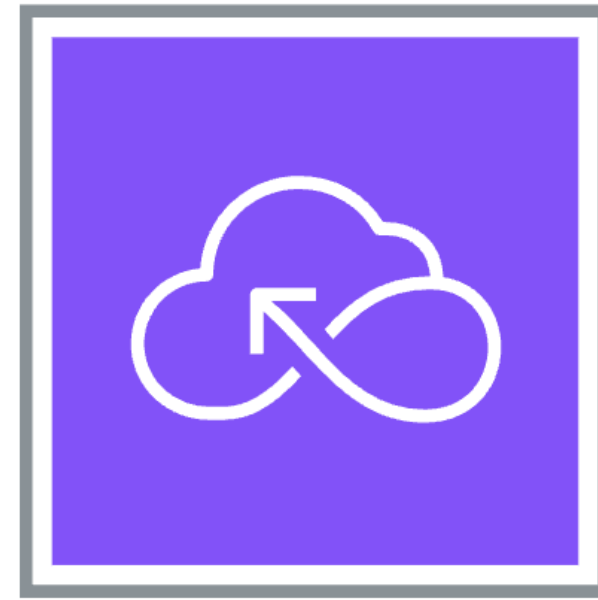


Automated infrastructure management

Containers and serverless compute



Containers



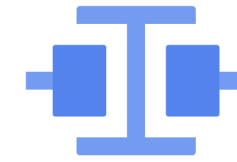
Serverless compute

What are containers?

Containers encapsulate applications and their dependencies, in lightweight singular units

Why containers?

- Isolate applications from underlying system dependencies




Isolation

What are containers?


Containers encapsulate applications and their dependencies, in lightweight singular units

Why containers?

- Isolate applications from underlying system dependencies
- Share host OS for efficient resource utilization



Isolation



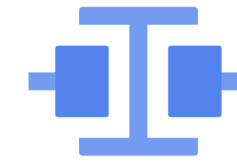
Increased efficiency

What are containers?


Containers encapsulate applications and their dependencies, in lightweight singular units

Why containers?

- Isolate applications from underlying system dependencies
- Share host OS for efficient resource utilization
- Easily movable and portable across environments



Isolation



Increased efficiency



Portability

Containers in AWS



Amazon Elastic Container
Service



Amazon Elastic Kubernetes
Service

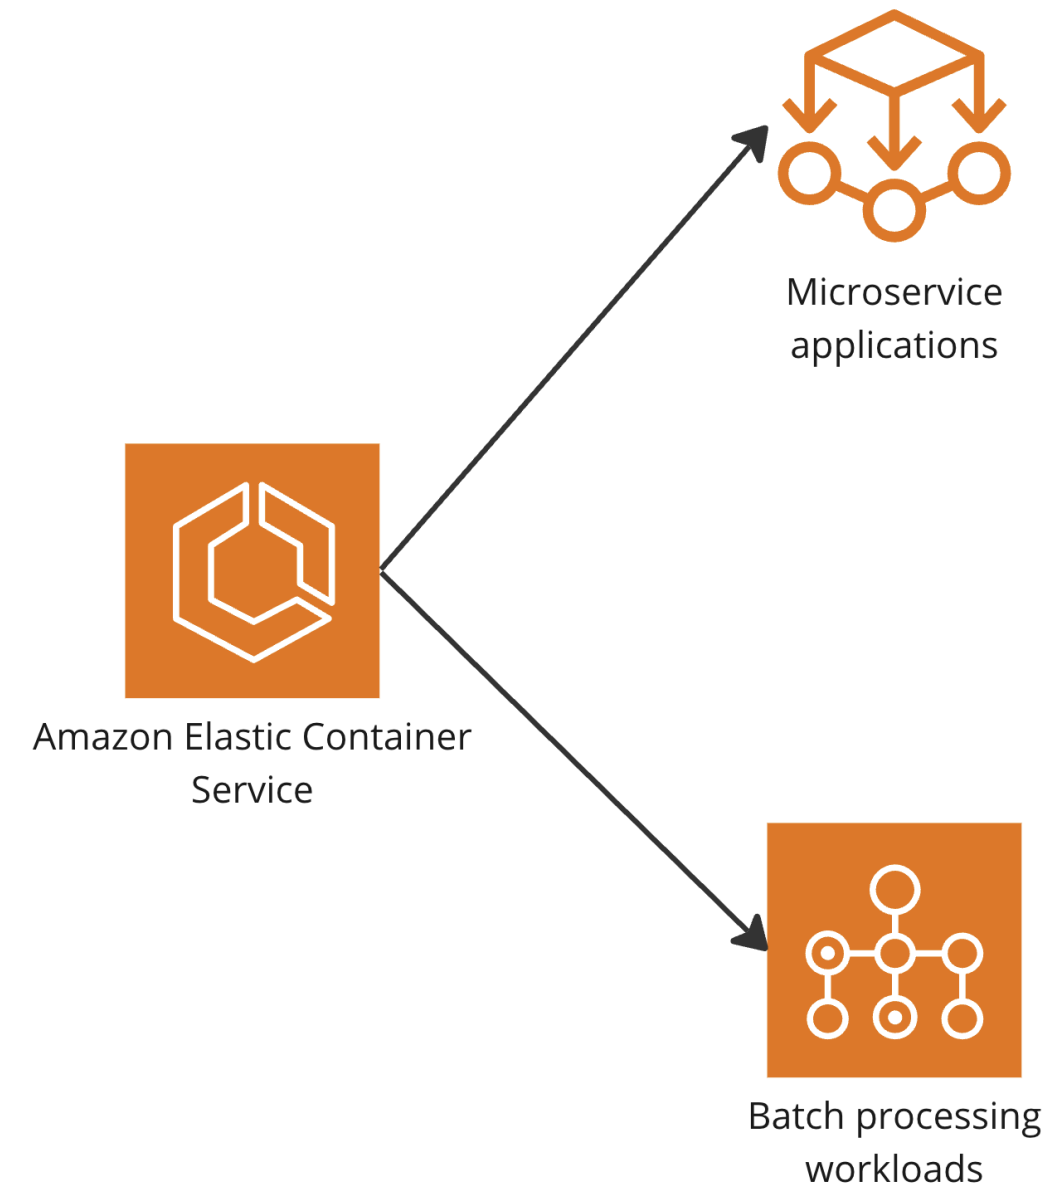
- Easily scale containerized applications up or down
- Integrate with other AWS services

Amazon ECS

Fully managed service for efficient deployment, management, and scaling of containerized applications

Use cases

- Deploying and managing microservices-based applications
- Plan, schedule, and run batch processing workloads across AWS services

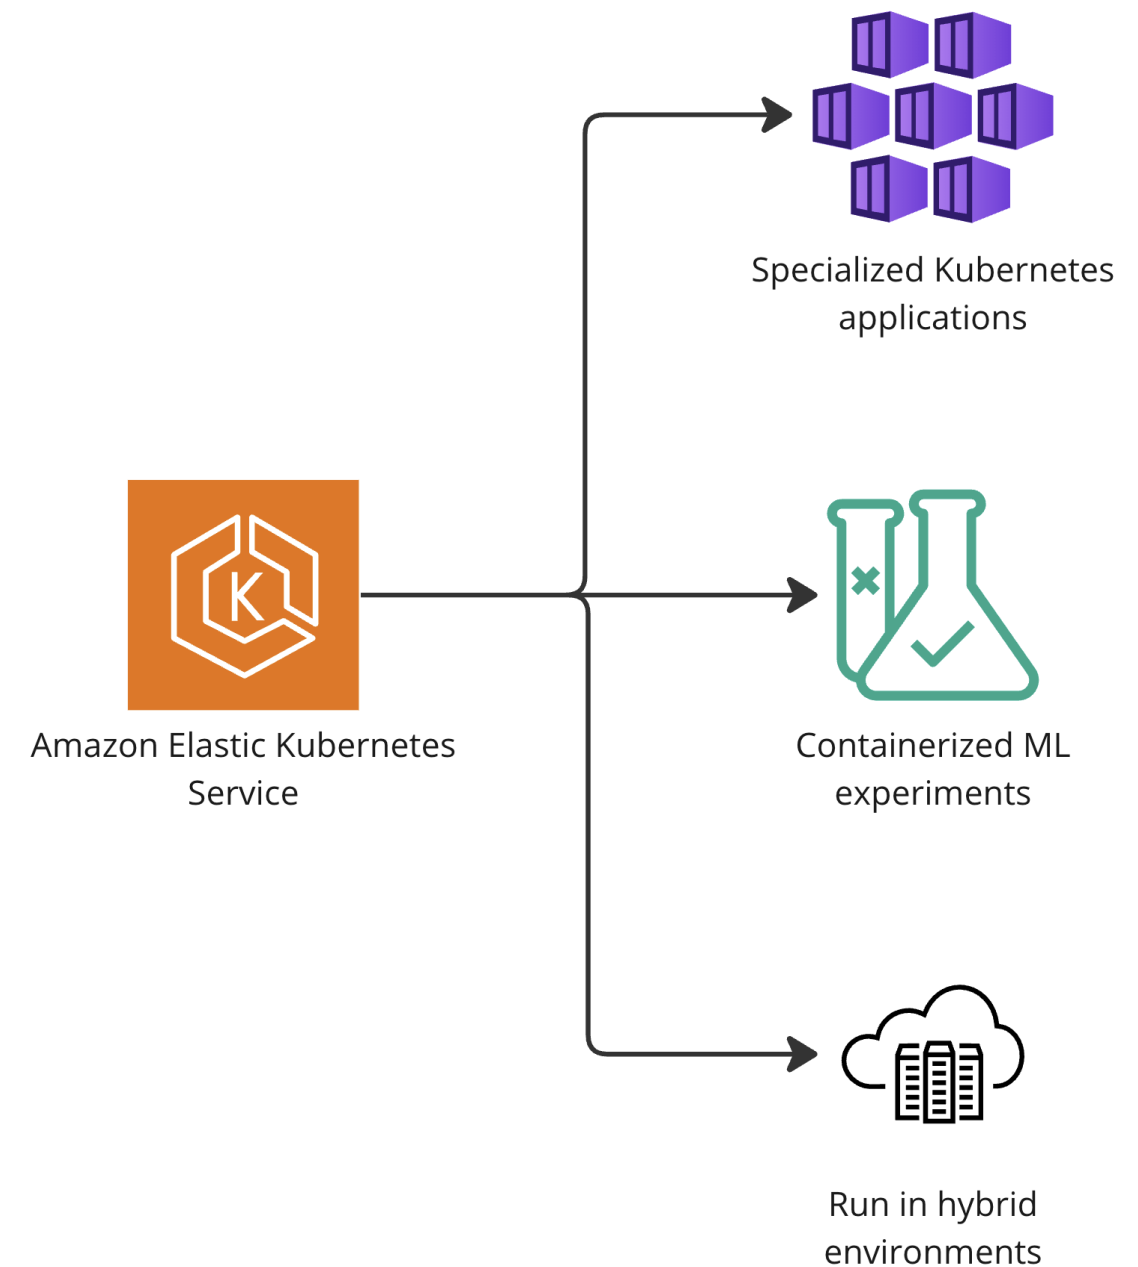


Amazon EKS

Container orchestration service specializing in running Kubernetes-powered applications

Use cases

- Pair with EC2 accelerated computing instances to run ML containers
- Manage clusters and applications in hybrid cloud environments



More forms of compute?

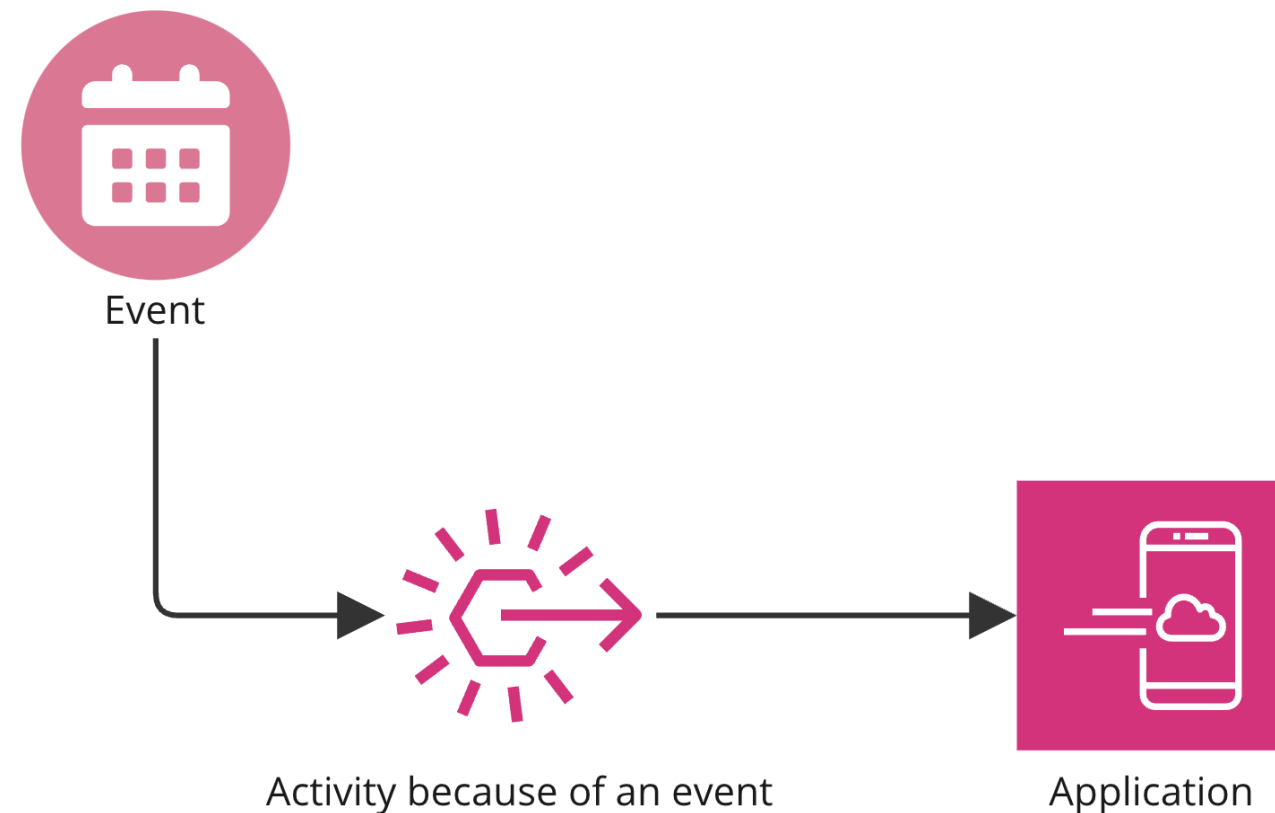


Containers

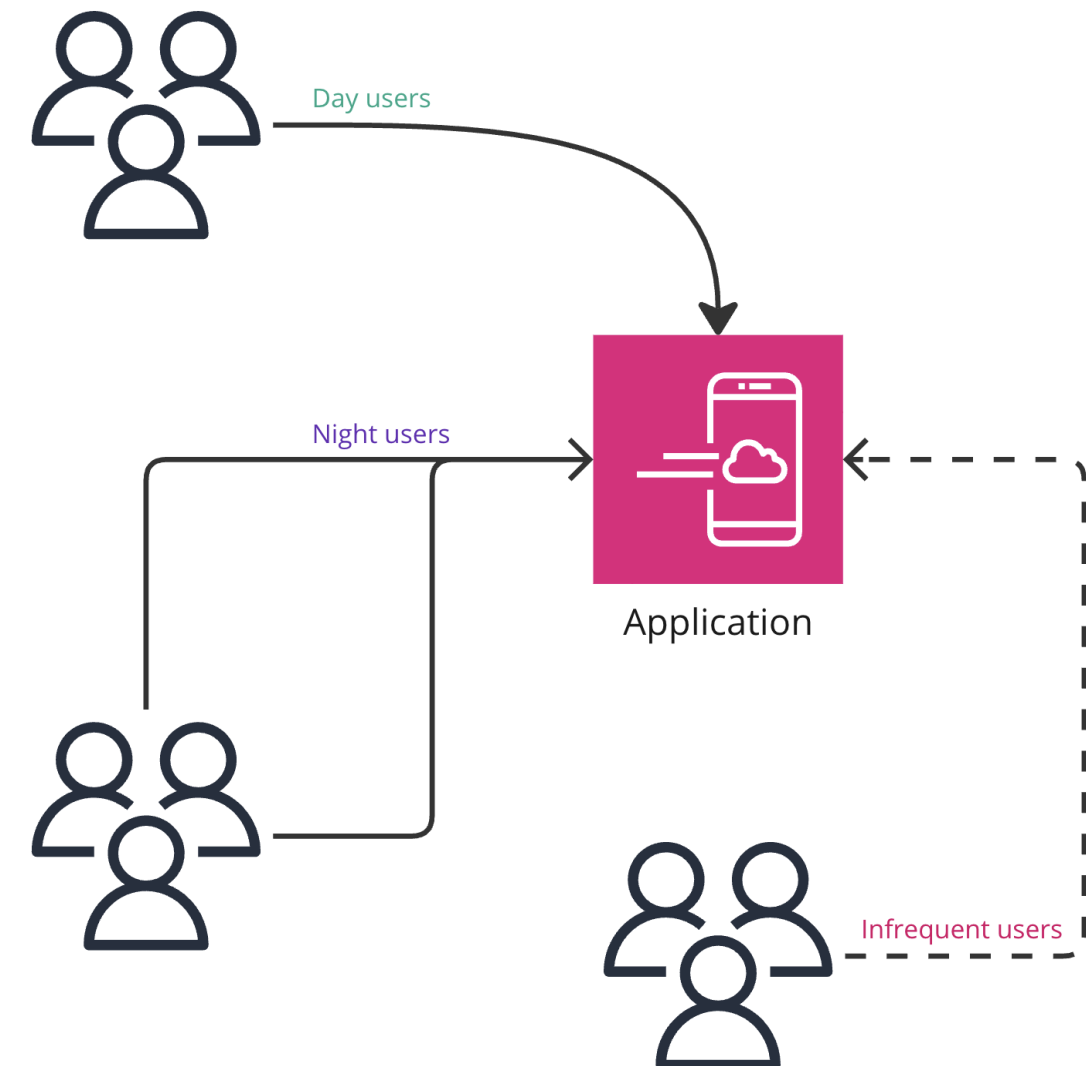
- Persistent environments
- Predictable workloads
- Resource-intensive applications

More forms of compute?

Event-based compute changes



Balancing compute for sporadic traffic loads



Serverless compute



Serverless compute

What is serverless architecture?

- No server management: forget about provisioning, scaling, or maintenance



No server
management

What is serverless architecture?

- No server management: forget about provisioning, scaling, or maintenance
- Event-driven: functions triggered by events in real-time



No server
management



Event-driven

What is serverless architecture?

- No server management: forget about provisioning, scaling, or maintenance
- Event-driven: functions triggered by events in real-time
- Cost-efficient: pay only for actual usage, not pre-allocated resources



No server
management

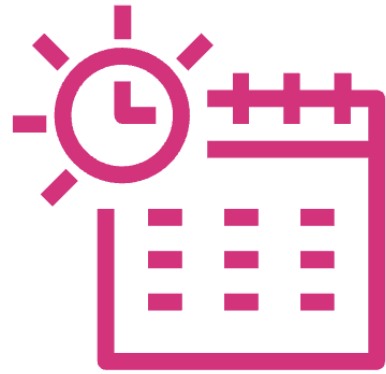


Event-driven



Cost-efficient

When to use serverless compute?




Event-driven applications



Real-time file processing



Uneven data bursts



Chatbots and voice assistants

Serverless compute in AWS



AWS Lambda



AWS Fargate

AWS Lambda

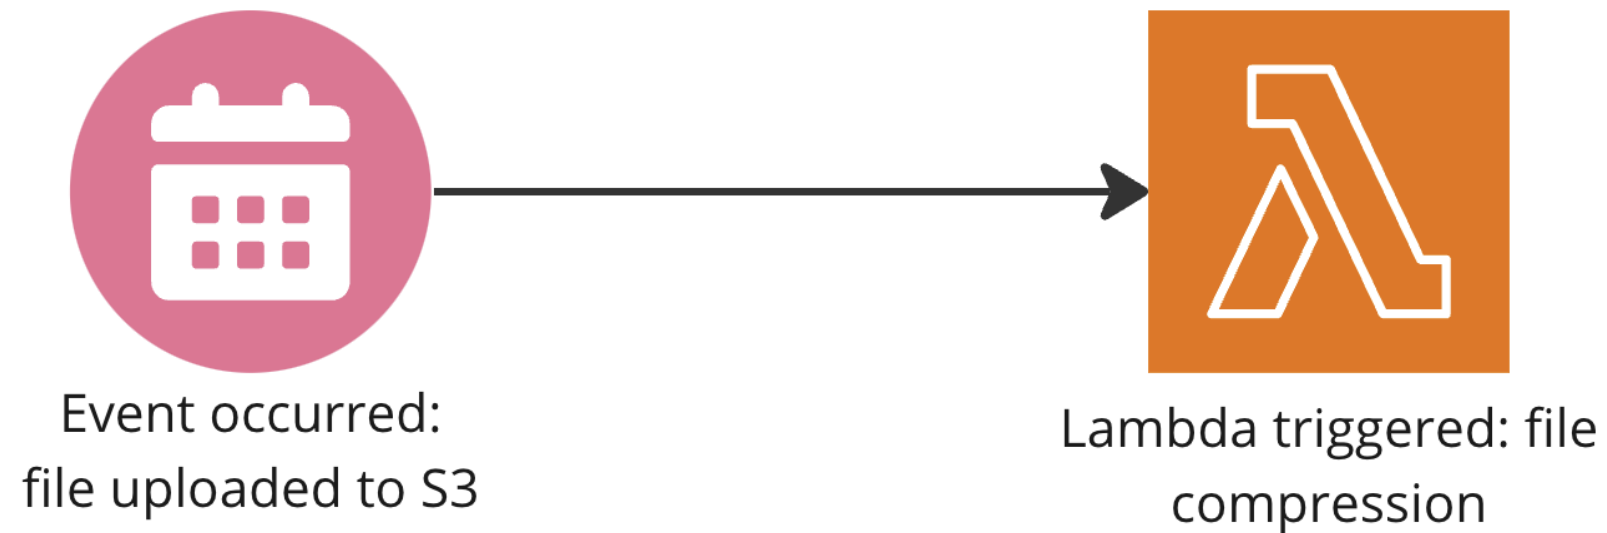
- Run code in response to events without provisioning or managing servers
- Automated compute scaling capabilities



Event occurred:
file uploaded to S3

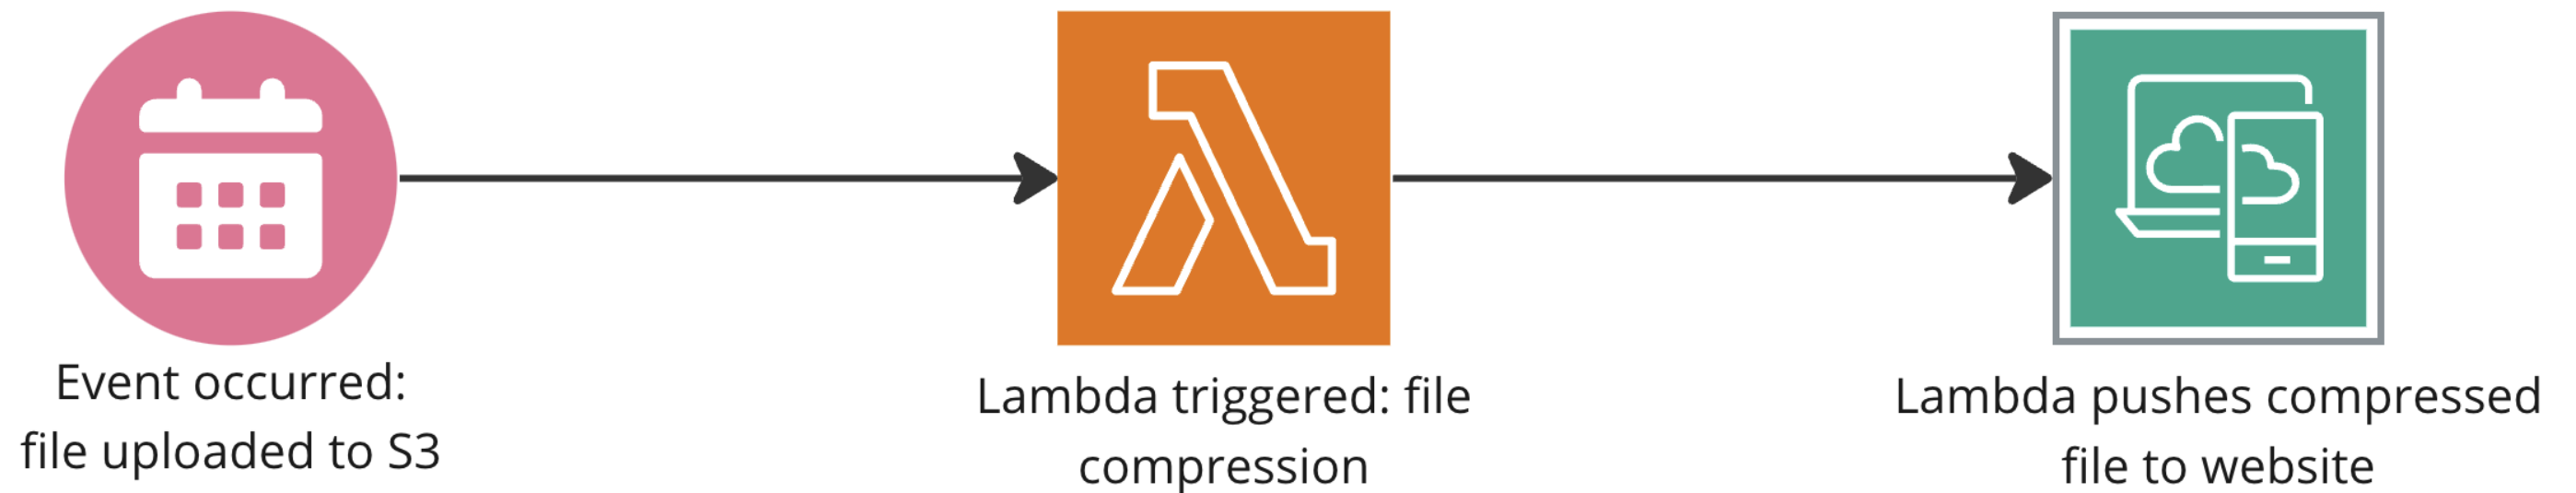
AWS Lambda

- Run code in response to events without provisioning or managing servers
- Automated compute scaling capabilities



AWS Lambda

- Run code in response to events without provisioning or managing servers
- Automated compute scaling capabilities

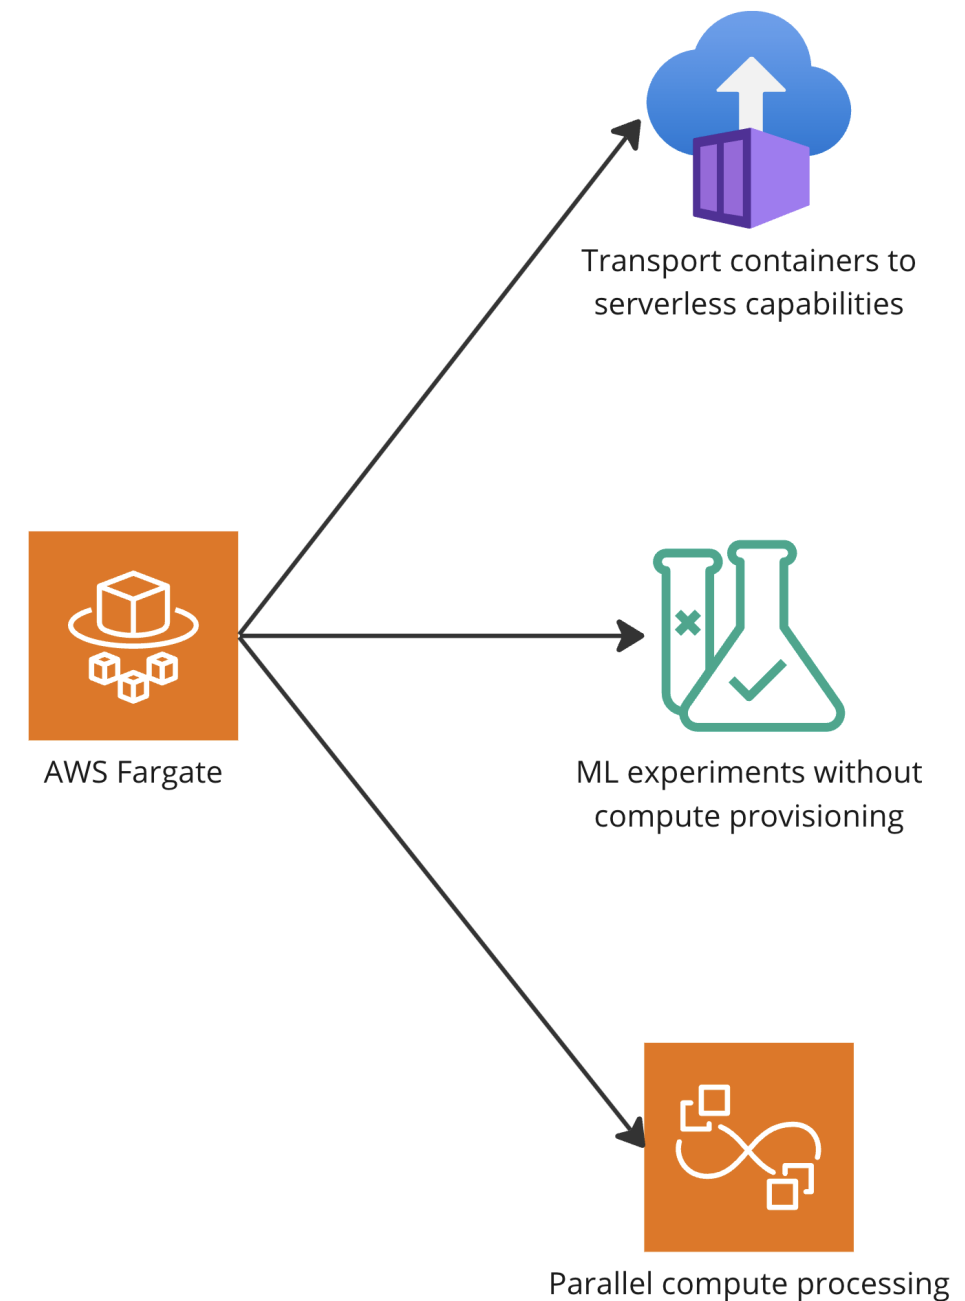


AWS Fargate

Streamlines application development by providing serverless compute for containers

Use cases

- Enable AI and ML applications without the need for excessive server provisioning
- Batch processing of large datasets with parallel compute capabilities



Let's practice!

AWS CLOUD TECHNOLOGY AND SERVICES CONCEPTS